

1 Data-driven recommendations for increasing reliability in measures of infant word recognition

2 Veronica Boyceⁱ, Virginia A. Marchman^a, Heidi Baumgartner^a, Claire Augusta Bergey^b,
3 Mika Braginsky^a, Federico Giovannetti^c, George Kachergis^a, Kyle MacDonald^j, Jess
4 Mankewitz^d, Stephan Meylan^e, Ben Prystawski^a, Robert Z. Sparks^a, Adrian Steffan^f, Alvin
5 Wei Ming Tan^a, Tian Linger Xu^g, Michael C. Frank^a, & Martin Zettersten^h

6 ^a Department of Psychology, Stanford University, Stanford, CA, USA

7 ^b Department of Linguistics, Stanford University, Stanford, CA, USA

8 ^c Unidad de Neurobiología Aplicada (UNA, CEMIC-CONICET), Buenos Aires, Argentina

9 ^d Department of Psychology, University of Wisconsin, Madison, WI, USA

10 ^e Department of Linguistics, University of California, Berkeley, CA, USA

11 ^f Department of Psychology, Ludwig Maximilian University, Munich, Germany

12 ^g Department of Psychology, Indiana University, Bloomington, IA, USA

13 ^h Department of Cognitive Science, University of California, San Diego, CA, USA

14 ⁱ Department of Brain and Cognitive Sciences, Massachusetts Institute of Technology,
15 Cambridge, MA, USA

16 ^j IBM

17 Author Note

18 AI Disclosure statement: Generative AI was used to do code refactoring and code
19 review.

Abstract

Children’s early word recognition is both an important phenomenon in its own right and a key index of early language more broadly. A key paradigm for measuring word recognition in young children is the looking-while-listening paradigm (LWL), which tracks eye movements to the referent of a word as a measure of online language comprehension. Yet the analytic and design choices for LWL vary widely, and the consequences of these choices have rarely been investigated at a large scale. Here we use data from Peekbank, a repository of LWL data from studies with almost 4000 children (6 months – 5 years), to search for analytic and design decisions that maximize the reliability and validity of LWL measures. Analytic recommendations include using long analytic windows for computing accuracy and avoiding baseline correction. We also find little or no benefits from excluding children or trials with low data. The recommended combination of analytic choices increased reliability (intra-class correlations and test-retest correlations) by approximately .20 and validity (correlations with receptive and productive vocabulary) by .07. Design recommendations include building experiments with a minimum of 10 target trials per child per condition and matching stimuli for animacy. We hope that these recommendations lead to improved and unified standards for the measurement of children’s word recognition in the LWL paradigm.

Keywords: looking while listening, word recognition, language development, eyetracking, reliability, validity

Word count: X

Data-driven recommendations for increasing reliability in measures of infant word recognition

Introduction

Over the first few months of life, children become attuned to their native language, organizing the phonological system of speech input and parsing speech into meaningful units. Soon after, children begin to show early signs of comprehension of words or phrases through simple behaviors such as reaching for an object that is named (e.g., “Where’s the ball?”). These everyday accomplishments mark important milestones in children’s development. Yet global assessments of these milestones do not capture the millisecond-level processes that take place during real-time language comprehension. Children, like adults, process speech incrementally, as it unfolds in time (Fernald, Pinto, Swingley, Weinberg, & McRoberts, 1998; Frank et al., 2025; Snedeker, 2013), and use features of the speech signal to extract meaning that support language comprehension (Law & Edwards, 2015; Weisleder & Fernald, 2013). These moment-to-moment processes are difficult to document without child-friendly procedures that precisely measure how children’s eye-movements are time-locked to speech. In particular, the “looking while listening” (LWL) procedure has emerged as a *de facto* standard for assessing language comprehension in infants and children (Frank, Braginsky, Yurovsky, & Marchman, 2021).¹ In the standard version of LWL, children see two pictures on a screen and hear speech directing their attention to one of the pictures, e.g., “look at the ball”. The accuracy and speed of fixating the named referent measure the efficiency of language comprehension (Fernald et al., 2008).

LWL has been a key tool for enhancing our understanding of the continuous processes of growth in incremental language processing (Golinkoff, Ma, Song, & Hirsh-Pasek, 2013). For example, using LWL methods, we can measure developmental increases in the speed and accuracy of children’s recognition of familiar words (Fernald, Perfors, & Marchman, 2006; Frank et al., 2025) and in their sensitivity to detailed phonological information during the first years after birth (Moore & Bergelson, 2024; Swingley & Aslin, 2000). Moreover, measures of children’s skill in the LWL task have proven useful as an index of individual differences. Children who show more efficient language processing of familiar words at 18 months show larger vocabulary size and more rapid vocabulary growth over the second and third years of life than children with less efficient processing (Fernald & Marchman, 2012; Fernald et al., 2006; Frank et al., 2025; Peter et al., 2019). Variation in language processing efficiency has also been shown to be associated with later language, cognitive, and non-verbal outcomes in both typically-developing (Law & Edwards, 2015; Peter et al., 2019) and preterm children (Loi, Marchman, Fernald, & Feldman, 2017; Marchman, Adams, Loi, Fernald, & Feldman, 2016; Marchman et al., 2018).

Despite widespread use of the LWL method in language development research, there is still limited consensus on how to best process and analyze time-series data of this type (from either adults or infants, Mirman, Dixon, & Magnuson, 2008; Bergelson & Swingley, 2012;

¹ Here we adopt LWL as a convenient label while recognizing that the procedure goes by a number of different names in the literature including the “intermodal preferential looking procedure” (Fernald, Zangl, Portillo, & Marchman, 2008; Hirsh-Pasek & Golinkoff, 1996; Reznick, 1990).

Peelle & Van Engen, 2021; Von Holzen & Mani, 2012; Weaver, Saffran, & Arunachalam, 2026). In the developmental field especially, this lack of consensus likely stems from structural factors: infant recruitment is time- and cost-intensive and datasets collected by individual labs are often small (Bergmann et al., 2018; Frank, Braginsky, Yurovsky, & Marchman, 2017). These issues make it difficult for any one researcher to systematically evaluate the consequences of specific analytic choices independently of the data collected for a particular study (cf. Peelle & Van Engen, 2021). Instead, researchers make ad-hoc decisions and adopt idiosyncratic lab practices that themselves originate from subjective judgments of what “seems to work”. By decreasing uniformity in the literature, this variation also increases the challenge of triangulating what precisely is being measured by LWL (Weaver et al., 2026). How can we put the measurement and analysis of infant word recognition and online language processing on firmer empirical footing?

Peekbank is a new data resource that offers a potential solution to this core issue by providing a large database of infant eye-tracking data from the LWL paradigm for exploring methodological questions (Zettersten et al., 2022). The most recent release of Peekbank (version 2026.1) houses 44 datasets from infants across a wide age range (6 months – 5 years) and across a broad set of test words and materials. The goal of the current manuscript is to use these data to systematically explore key analytic decisions that researchers make when using this paradigm, offering best-practices guidance for measurement.

Our approach throughout is to seek analytic choices that distill the best possible measures from the rich time course of behavior in the LWL task. Recent reviews of the LWL literature highlight open questions about both the reliability and validity of looking-based measures of word recognition (Weaver et al., 2026). Starting with reliability, we aim for analyses that yield the maximal signal about individual children relative to both measurement error and irrelevant artifacts (Byers-Heinlein, Bergmann, & Savalei, 2022; Rhemtulla & Bottesini, 2022; Schreiner et al., 2024). Here, to measure reliability, we compute both within-dataset inter-item reliability and test-retest reliability for datasets with multiple LWL sessions close in time for the same children.

Turning to validity, the goal is to make analytic decisions that maximize the relationship with other measures that are thought to reflect related constructs (Flake & Fried, 2020; Kominsky, Lucca, Thomas, Frank, & Hamlin, 2022). Here, we select as our primary measure the correlations between LWL measures and parent reports of vocabulary measured with aggregate measures of receptive and expressive vocabulary size from the MacArthur-Bates Communicative Development Inventories (MB-CDI), a popular parent report measure of early vocabulary (Frank et al., 2021; Marchman, Dale, & Fenson, 2023) that is widely used in conjunction with LWL tasks (CITE Fernald etc).

Our focus in this paper is to provide guidance for navigating some of the many decisions facing an investigator analyzing LWL data. Previous work has provided in-depth recommendations for experiment design, eye-tracking, and eye-gaze coding (Fernald et al., 2008; Olson et al., 2020; Venker et al., 2020; Wass, Smith, & Johnson, 2013). Other recent work using an earlier version of Peekbank (as well as other datasets) investigated the

psychometric properties of a range of different dependent measures in LWL tasks and their general construct structure (Sander-Montant, Fibla, Byers-Heinlein, & Killam, 2025). We focus here on the consequences of varying analytic choices for two main measures which tend to have the strongest psychometric properties (Sander-Montant et al., 2025). The first is *accuracy*: proportion looking to the target image (out of all looks to either the target or distractor image) computed over some period of time after the onset of the target word in the auditory stimulus. The second measure is *reaction time* (RT): how long it takes a child to shift their gaze from the distractor to the target image in response to the auditory stimulus (Fernald et al., 2008). For both measures, the goal of data processing should be to include the timepoints, trials, and sessions that capture signal that could represent word knowledge and language processing, while removing noise. While examining the impacts of different potential data processing choices, we identified several experimental design questions that had notable consequences for our central measures.

The plan of the paper is as follows. We begin by overviewsing the subset of the Peekbank database we selected for our analyses. We then make a set of targeted recommendations, providing evidence for each based on the impact of our recommendation on the average reliability and validity of the resulting measures. Our first set of recommendations focus on data processing decisions and our second set of recommendations concern design considerations. To foreshadow our results, we find that while many common data analysis practices provide relatively reasonable levels of reliability and validity, it may be possible to further optimize reliability and validity through careful, data-driven decision-making.

Methods: Dataset and approach

Dataset

Peekbank is an aggregation of datasets from a number of LWL studies. This work is based on Peekbank version 2026.1 (released February 2026) which houses 44 datasets containing 134929 trials from 3915 children participating in 1040 distinct testing sessions (which we term “administrations”; some children participate in multiple sessions).

We limit our analyses to what we term “standard” trials, resulting in 75636 trials from 3171 children participating in 4851 sessions, across 36 datasets. For a trial to be considered “standard”, both the target and the distractor images must be of familiar objects. For the auditory stimulus, the target word must be a real word that occurs once and is the first and only source of information about the target. The carrier phrase must be a well-formed, grammatical phrase in a single language (e.g., English, Spanish), and without any intentionally-added noise or audio filtering. Standard trials could come from studies in which this was the only trial type tested or from studies in which familiar word trials were used as fillers or baseline trials presented in conjunction with other trial types. (See SI: Criteria for Inclusion for more detailed criteria used to handle borderline situations.)

Following Peekbank’s convention, stimulus onset is defined as the onset of the target word and labeled as the zero point for timing. We only include trials from children 60

months or younger (see Figure S1 for age distributions). The average age in the final dataset was XYZ (SD; range; any other age infoTODO).

MB-CDI comprehension scores were available for 632 children from 12 of the datasets, and MB-CDI production scores were available for 1143 children from 20 of the datasets. Test-retest data comes from 5 datasets with multiple sessions within a month of each other, covering 365 children (788 session pairs). See Table 1 for a list of included datasets. Figure S2 has time course plots showing the number of trials for each dataset.

Table 1

Descriptives of the included datasets.

	Dataset Name	Pct Trials	Trials	Admins	Subjects	Age	CDI	Test-retest	Language
1	Adams et al. (2018)	18.4%	13899	711	69	22 [13 - 38]	x	x	English
2	Fernald & Marchman (2012)	15.6%	11828	679	122	22 [17 - 32]	x	x	English
3	Weaver et al. (2024)	6.5%	4948	248	141	16 [14 - 24]		x	English
4	Fernald et al. (2013)	5.8%	4359	178	80	20 [17 - 26]	x	x	English
5	Hurtado et al. (2008)	5.3%	4035	136	76	21 [17 - 27]	x		Spanish
6	Fernald et al. (2006)	4.9%	3692	229	63	20 [15 - 25]	x		English
7	Yurovsky et al. (2013)	3.7%	2810	413	413	34 [12 - 60]			English
8	Yurovsky et al. (2017)	3.1%	2350	327	327	37 [8 - 60]			English
9	Bergelson & Swingley (2012)	2.8%	2081	84	84	12 [6 - 21]	x		English
10	Kartushina & Mayor (2019)	2.7%	2072	69	69	7 [6 - 9]	x		Norwegian
11	Weaver & Saffran (2026)	2.4%	1788	69	69	19 [18 - 20]	x		English
12	Weisleder & Fernald (2013)	2.3%	1769	58	29	22 [18 - 27]			Spanish
13	Yurovsky & Frank (2017)	2.3%	1747	285	285	26 [13 - 59]			English
14	Borovsky & Peters (2019)	2.2%	1676	79	79	18 [17 - 20]	x		English
15	Potter & Lew Williams (2023)	2.1%	1626	67	67	24 [21 - 27]	x		English
16	Yoon et al. (2015)	2.0%	1500	199	199	43 [13 - 60]			English
17	Sander-Montant et al. (2023)	1.8%	1381	123	123	22 [12 - 31]	x		English, French
18	Mahr et al. (2015)	1.6%	1232	29	29	21 [18 - 24]	x		English
19	Egger et al. (2020)	1.4%	1052	54	54	18 [18 - 19]	x		Dutch
20	Garrison et al. (2020)	1.4%	1046	35	35	15 [12 - 18]	x		English
21	Hurtado et al. (2007)	1.3%	1010	49	49	24 [15 - 37]			Spanish
22	Pomper & Saffran (2017)	1.0%	785	82	82	17 [14 - 19]			English
23	Frank et al. (2016)	1.0%	781	108	108	35 [12 - 60]			English
24	Bacon & Saffran (2022)	1.0%	754	38	38	23 [22 - 24]	x		English
25	Ronfard et al. (2022)	1.0%	737	40	40	20 [18 - 24]	x		English
26	Perry & Saffan (2017)	0.9%	718	42	42	20 [19 - 22]	x		English
27	Casillas et al. (2017)	0.9%	651	23	23	32 [9 - 48]			Tseltal
28	Swingley & Aslin (2002)	0.8%	596	50	50	15 [14 - 16]	x		English
29	Pomper & Saffran (2016)	0.6%	471	60	60	44 [41 - 47]			English
30	Kremin et al. (2023)	0.5%	410	60	60	42 [37 - 49]			English, French, Spanish
31	Pomper & Saffran (2015)	0.5%	389	25	25	40 [38 - 42]			English
32	Moore & Bergelson (2022)	0.5%	384	29	29	18 [16 - 20]	x		English
33	Potter et al. (2019)	0.4%	309	44	23	23 [18 - 29]	x	x	English
34	Byers-Heinlein et al. (2017)	0.4%	284	48	48	20 [19 - 21]	x		English, French
35	Pomper & Saffran (2019)	0.3%	249	44	44	40 [38 - 43]			English
36	Pomper & Saffran (2018)	0.3%	217	37	37	39 [38 - 43]			English
Total		100%	75636	4851	3171	22 [6 - 60]	21	5	

Approach

Reliability refers to how consistently measurements are measuring the same thing. One measure of reliability is an intra-class correlation (ICC) that measures (within a dataset) how much different items (target labels) are consistent with each other in “rating” the subjects (kid-administrations). See SI: Details about ICC for more specifics. Another measure of reliability is test-retest reliability, which measures how similarly the same child performs when assessed on two different occasions on the same task. This measure is only possible on a subset of the datasets that repeated administrations at closely spaced intervals. While test-retest reliability data are limited, this is the metric that is most important for measuring the stability of individual differences.

Validity refers to whether measures correspond to the construct of interest (in this case, language processing efficiency and/or receptive/expressive vocabulary size). Our validity measure is the correlation between RT and accuracy and CDI production and/or CDI comprehension. Since some studies used different versions of the CDIs (e.g., short vs. long forms), all scores were converted to a proportion of the possible items. In general, correlating with production and comprehension in the expected directions is desirable, although we also may expect that LWL measures are detecting distinct aspects of language learning compared to what is represented in the CDI (see Frank et al., 2025 for a detailed analysis).

All the code and data to reproduce the analyses is available at <https://github.com/peekbank/peekbank-method> [TODO: anonymize for review if needed].

Data processing recommendations

Recommendation 1: Measure accuracy over a long time window.

Accuracy, or proportion of target looking, must be computed over a time window of interest. Typically, windows start shortly after the onset of the target word and continue until a few seconds later; this window is intended to capture the peak of children’s stimulus-driven looks to the target. However, in practice, studies vary in what windows are used, and often window selection is a locus for analytic flexibility (choosing windows post hoc to maximize a particular effect) (Peelle & Van Engen, 2021). Common analysis windows include starting 300 milliseconds after noun onset and ending 1800 ms after noun onset (following Fernald et al., 2008) or 367ms to 2000 ms after stimulus onset (e.g., Swingley & Aslin, 2002), although windows that extend until 3 or 4 seconds after noun onset are also used (e.g., Bergelson & Swingley, 2012; Bergelson & Aslin, 2017; López Pérez, Moore, Sander-Montant, & Byers-Heinlein, 2024). In this analysis, we begin from first principles and ask when windows should start and end to maximize reliability and validity on word recognition trials.

Analytic approach. We first calculated the ICC reliability for accuracy across datasets, computing average ICC across a range of different window onsets and offsets. Next, we added the additional metrics of test-retest reliability and—to assess validity—CDI correlations, which were calculated by dataset for each combination of parameter values. To aggregate across datasets, we used a linear mixed-effects model with random intercepts by dataset, weighted by the number of administrations. We interpret the fixed-effect contrasts as what differences in outcome measure would be expected from switching from one set of parameters to another on a generic dataset. We report estimates and 95% confidence intervals. The best data-processing options might be different for children of different ages, so we ran the same linear mixed effect models separately for data just on children in each of 4 age groups: <18 months, 18-24 months, 24-36 months, and >36 months.

Results. Average ICCs across windows are shown in Figure 1. As expected, including time points from before the noun onset decreases reliability. Our initial exploration found a region of high reliability with onsets around 600-800ms and offsets around 2800-4000ms. A dataset by dataset comparison of the differences between a 300-1800ms window (representative of common practice) and a 600-4000 ms window are shown in Figure 2. Nearly all datasets have higher ICC reliability for the later and longer window. All 5




Figure 1. A heat map showing the average ICC across datasets for different start and end points for the accuracy window. A region with the highest ICC (best reliability) is outlined in black, and the region most commonly represented in the literature is outlined in blue.




Figure 2. Per dataset outcome measures for two window options: 300ms-1800ms (red) and 600ms-4000ms (blue). Dot size is proportional within each panel to the number of administrations.

test-retest reliability datasets also show increased reliability on this later and longer window. Our mixed effect model confirms that there are expected reliability gains from switching from a 300-1800ms window to a 600-4000ms one. ICC reliability is expected to increase by 0.12 [0.08, 0.16] and test-retest reliability by 0.07 [-0.02, 0.16]. There is no evidence of differences in validity between these two windows (-0.02 [-0.07, 0.03] for vocabulary comprehension or 0.03 [-0.02, 0.08] for vocabulary production).




Figure 3. Per dataset ICC measures for two window options: 300ms-1800ms (red) and 600ms-4000ms (blue), shown split by age category. Dot size is proportional within each panel to the number of administrations.

These patterns generally hold across age. The overall level of reliability rises with age, as older children generally fixate target images more consistently and experiments with younger children typically use fewer trials (Figure S3). However, the overall pattern of which time windows maximize reliability and validity at any given age is remarkably consistent across the different age groups (Figure S5). For instance, across different ages there is an improvement in ICC for using a 600-4000 ms window compared to a 300-1800 ms one (Figure 3). For <18 months, the ICC difference is 0.06 [-0.03, 0.16]; for 18-24 months, 0.13 [0.07, 0.19]; for 24-36 months, 0.06 [0.01, 0.12]; and for 36+ months 0.15 [0.06, 0.23].

While we focused our analyses on the window that maximized reliability numerically (600 - 4000 ms), there are a wide range of window options with generally comparable reliability and validity measures (Figure S4). For instance, a 400 - 3000 ms window is not significantly different from a 600 - 4000 ms window on any of the outcome measures. Although we recommend starting between 400ms and 600 milliseconds to optimally balance validity and reliability, choosing when precisely the accuracy window starts is likely less important for maximizing reliability than choosing a long time window: regardless of window start time, increasing the window size consistently improves reliability (Figure 1).

Discussion. Overall, we find that LWL accuracy measures are reliable across a wide range of window choices; however, to maximize reliability, we recommend windows that start somewhere in the 400-600 ms range and extend to 3000-4000 ms (longer than the canonical window that ends at 1800 or 2000 ms). Windowing choices seem to primarily affect reliability measures and have limited effects on our validity measures. The finding that later starting and later ending windows increase reliability suggests that, even as looks target decline in the post-2000ms window, looking behavior still contains reliable signal. Conversely, there may be relatively little reliable signal in the early 300-600ms looking period. While

some children may look quickly to the target and then look away, other children who look to the target later or who continue looking at the target are providing a useful signal during later portions of a given trial. Overall, children in many datasets are in aggregate looking to the target more than the distractor during later looking periods (even if less than at the peak, Figure S2).

Recommendation 2: Do not use baseline-corrected accuracy.

One commonly acknowledged issue in LWL studies is that looking behavior could be driven by differences in visual salience between the target and distractor images themselves (Fernald et al., 2008; Novack et al., 2025): Children may have a baseline preference to look at one image over another, even in the absence of auditory stimuli. Although thoughtful design choices can minimize these issues (see Recommendation 6), some analytic strategies also attempt to account for baseline preferences in dependent measures. While there are a variety of such strategies in the LWL literature, here we focus on one common practice: baseline correction. In baseline correction, researchers use a difference measure between accuracy on a critical post-stimulus window and “accuracy” (proportion looks to target among looks to target or distractor) on a pre-stimulus window (e.g., Bergelson & Swingley, 2012; Weaver, Zettersten, & Saffran, 2024). The pre-stimulus window can be the 2 seconds prior to disambiguation (Weaver et al., 2024) or the full 3-4 second pre-stimulus period (Bergelson & Swingley, 2012).

Analytic approach. We considered baseline-corrected accuracies with various pre-stimulus windows, focusing on a 600-4000 ms critical window (although the results generalize to other critical windows). We restricted analysis to datasets with at least some data from 2000 ms or more before noun onset (see Figure S2 for which datasets have data when). We also computed a “no baseline” accuracy, based solely on accuracy within a 600-4000 ms time window with no correction. Baseline windows started at -1000, -2000, -3000 and -4000 ms and ended at noun onset.

Results. Baseline-corrected accuracy led to worse reliability and validity than accuracy. As an illustrative example, Figure 4 shows per-dataset results for accuracy compared to a baseline-corrected accuracy using a -2000 - 0 ms baseline window. This result held generally: All baseline window sizes that we considered led to worse reliability and validity measures compared to using normal accuracy without any correction. For a regular window of 600-4000 ms, baseline correction is expected to decrease ICC by at least 0.1 [0.04, 0.16]; test-retest reliability by at least 0.19 [0.08, 0.3]; correlation with receptive vocabulary by at least 0.04 [-0.05, 0.12] and correlation with productive vocabulary by at least 0.09 [0.05, 0.13]. Similar patterns hold for different accuracy windows (Figure S6) and across different age groups (Figure S8).

Discussion. We do not recommend the use of baseline-correction in accuracy measures, as it worsens both reliability and validity compared to a plain accuracy measure. Why would incorporating baseline information lead to worse metrics? The literature on the psychometrics of difference scores (e.g., Lord, 1956; Rogosa & Willett, 1983; Trafimow, 2015) provides a valuable guide for evaluating the expected measurement properties of baseline correction. This literature points to two important features of LWL tasks: first, children’s looking behavior is noisy from trial to trial, meaning that the reliability of proportion




Figure 4. Comparison between a baseline-corrected accuracy in red and accuracy in blue across datasets and measures. Both use an accuracy window of 600 - 4000ms, the baseline window was -2000 - 0ms.

looking during both the baseline and the critical window is low; and second, the signal across baseline and critical portions of a trial is typically correlated (in part due to the very salience that the baseline correction is trying to remove). Unfortunately, it is under precisely these conditions – low reliability and high correlation between components – that difference scores exhibit the lowest reliability (Trafimow, 2015). Thus, while baseline-correction may lead to grand average accuracy scores that are less biased (by removing item salience), it also leads to less reliable and valid individual-level measures.

Recommendation 3: Measure land-based reaction times with a minimum reaction time of around 400 ms.

Reaction time (RT) reflects the speed with which infants shift their gaze from the distractor to the target image. RT computations include many choice points. RTs can be computed based on when a child's eyes leave from the distractor (launch-based RT) or when a child's eyes arrive on the target image (land-based RT) (Fernald et al., 2008). A minimum and maximum RT may be set, as short and long RTs are often excluded because they are considered unlikely to be generated in response to the stimulus (Fernald et al., 2008). Reaction time data are sometimes analyzed using a logarithmic transform (consistent with the idea that reaction times are log-normally distributed and that effects on them may be multiplicative rather than additive) (Csibra, Hernik, Mascaro, Tatone, & Lengyel, 2016; Frank et al., 2025).

Analytic approach. We considered a wide range of options for measuring RT, fulling crossing a) launch- or land-based RTs, b) a minimum RT threshold for inclusion between 0 and 1000 ms, c) a maximum RT threshold for inclusion between 1600 and 4000, and d) raw or log-transformed RTs.




Figure 5. Heat map of different parameter options for measuring RT and how they on average affected ICC. A region with the highest ICC (best reliability) is outlined in black, and the region most commonly represented in the literature is outlined in blue.




Figure 6. Comparison between two different sets of RT parameters. In red, an example of common practice, using a launch-based RT from 300-1800ms. In blue, an RT that maximizes ICC, with a land RT from 400 to 4000ms, on the log scale. Note that correlations with the CDI are generally expected to be negative (as faster word recognition is associated with larger vocabulary).

Results. An initial exploration of ICC reliability (Figure 5) indicated that ICC reliability was maximized for land-based RTs; a minimum threshold for including RTs set at around 400 ms; and with higher maximum thresholds for RT inclusion. A comparison between a common set of parameters for calculating RT and an alternative that maximizes ICC are shown in Figure 6. Compared to computing the “standard” launch-based, raw RTs restricted to 300-1800ms, using land-based log RTs ranging between 400-4000ms is expected

to increase ICC by 0.11 [-0.01, 0.23], with little change in test-retest reliability (0.04 [-0.18, 0.27]). These choices had little impact on CDI correlations (-0.04 [-0.17, 0.08] for comprehension, and -0.03 [-0.12, 0.06] for production). Similar patterns hold for children of different ages (Figure S9 and Figure S11).

Discussion. Overall, RT measures capture substantial signal, as indicated by the moderate ICC values across all choice parameters. A wide range of parameter options have similar properties (Figure S10 and Figure S12), but computing land-based RTs with an inclusion window of 400 - 4000ms maximizes reliability without negatively impacting validity. Note that the recommendation to compute land-based reaction times contradicts traditional approaches to reaction-time computation (Fernald et al., 2008) based on launch time; however, land-based RTs generally led to higher ICCs across our dataset. Log-scale transformations should be considered, but may depend on modelling and theoretical considerations.

Recommendation 4: Trial-level and session-level exclusions for low amounts of data are generally unnecessary.

A common issue in developmental research is when to apply trial- and session-level exclusions. For LWL studies, trial-level exclusions tend to be for contributing too few time points (e.g., due to tracker loss or looking away from the screen during the trial). Papers often use a combination of heuristic exclusions for inattention prior to coding looking behavior (Adams et al., 2018) and criterion-based exclusions trials based on the proportion of looking to the target and distractor during the trial. Cutoffs vary, with different experiments requiring minimum looking at the images of at least 20%, 50% or 66% of the critical window of a trial (Borovsky & Peters, 2019; Pomper & Saffran, 2019; Weaver et al., 2024). Studies also occasionally filter for looking at the target or distractor specifically at the point of disambiguation (C. E. Potter, Fourakis, Morin-Lessard, Byers-Heinlein, & Lew-Williams, 2019).

In addition, studies also implement session-level exclusions based on the minimum number of valid trials. For accuracy measures, children are often excluded if they had fewer than 25% or 50% of the total trials in a session, either within a condition, or overall (Adams et al., 2018; Garrison, Baudet, Breifeld, Aberman, & Bergelson, 2020; Perry & Saffran, 2017). For RT, it is instead standard to require a minimum number of trials [e.g., two; Fernald et al. (2006); Adams et al. (2018)].

Analytic approach. We consider 3 types of exclusions: trial-level exclusions for accuracy, administration-level exclusions for accuracy, and administration-level exclusions for RT. We first explored trial-level accuracy exclusions, both on the basis of proportion of off-screen looks and on the basis of where the child was looking at noun onset. For accuracy, we then considered per-session minimums either based on the number of valid trials, or based on the fraction of valid trials in an administration. The fraction of valid trials is the more common exclusion in the literature, but when drawing from many datasets that vary in the total number of trials, minimum trial count may represent a more comparable metric. For RT, we considered different cutoffs for the minimum number of valid trials required for inclusion.

We note that datasets contributed to Peekbank vary in whether excluded data was contributed, or whether only cleaned post-exclusion data was contributed. Thus, some datasets may not show effects of exclusions if those trials or administrations that would be excluded are missing from Peekbank.

Results. For exclusions, there are both considerations regarding reliability and validity, and regarding how much of the potentially usable data is being discarded. For trial-level exclusions, a requirement to look at the target and distractor at least 50% of the time results in 85.64 % (range: 27.22% - 98.92%) of trials being retained. These exclusions lead to an expected increase in ICC by 0.03 [0, 0.05] and no substantial change in test-retest reliability (-0.02 [-0.11, 0.08]). Differences in CDI correlation are not substantive (0.01 [-0.04, 0.06] for comprehension and 0 [-0.03, 0.03] for production).

Filtering based on initial looking position, or having requirements for looking above 50% start to diminish the amount of data, and potentially negatively affect test-retest reliability and CDI measures (Figure S13 and Figure S14). Requirements for on-screen looking rates between 0% (i.e., no exclusions) and 50% lead to globally similar outcomes numerically.

We then explored the effects of excluding sessions for too few trials. Here, we focus on the results for a minimum number of 8 total trials, with no trial-level exclusions, to illustrate general patterns in the results, but report outcomes for a range of proportion- and count-based thresholds in the supplement (Figures S15-S18). Only including sessions with at least 8 trials results in retaining 80.99 % (range: 0% - 100%) of total sessions. These exclusions show little evidence of increasing ICC (0 [-0.01, 0.02]) and or test-retest reliability (0.03 [-0.08, 0.14]) compared to retaining all sessions. Differences in CDI correlation are not substantive (0.01 [-0.04, 0.07] for comprehension and 0.01 [-0.03, 0.05] for production). We do not observe significant differences based on minimum trials below 10 or minimum fractions (Figures S15-S18). Overall, excluding sessions with too few trials had minimal impact on reliability and validity.

We also considered the effects of minimum trial requirements on RT data. Only including sessions that contribute at least 2 trials results in retaining 90.93 % (range: 42.86% - 100%) of all sessions. Similar to previous results, these exclusions lead to a negligible change in ICC by -0.01 [-0.07, 0.04], a slight numerical increase in test-retest reliability by 0.03 [-0.08, 0.14], no change in CDI comprehension correlation (0 [-0.15, 0.14]), and a slight numerical increase in CDI production correlation (-0.02 [-0.08, 0.04]). For RT, setting a low minimum threshold for number of trials (e.g., 2 trials) appears to lead to moderate improvements in test-retest reliability and slightly stronger correlations with CDI production (Figures S19 and S20).

Discussion. Overall, exclusions based on amount of data per-trial or per-session come with trade-offs between the amount of data and variability within the data. Standard practice within the field often errs on the side of excluding a large number of trials and sessions for missing data. Surprisingly, our results show that applying no exclusions or minimal exclusions often represents a reasonable choice, at least in terms of the psychometric properties of the resulting measurements.

Design recommendations

Our focus for this work was to explore how data processing choices affected the reliability and validity of the resultant measures. However, along the way, we identified several experimental design questions that had notable consequences for our central measures. We summarize two design recommendations that stood out as having the largest potential impact on reliability and validity (see supplement, section S8 for an additional analysis of order effects).

Recommendation 5: Design studies to have a minimum of 10 trials per child, per condition, or as many trials as possible for individual-difference studies.




Figure 7. Simulations of measurement properties for different trial counts, created by downsampling trials for all children with different trial minimums. Shown for an accuracy window of 600 to 4000 ms with no trial- or administration- level exclusions.

We found limited evidence for excluding children for low trial counts, but intuitively, more repeated measures should improve measurement properties such as reliability and validity. Rather than retrospective exclusions, another approach is prospective: how many trials per child should one aim for to achieve adequate reliability and validity? Test-retest reliability is important for studies that investigate longitudinal effects, as it assumes that the test sessions are representative of properties of the child at that age (Byers-Heinlein et al., 2022). To address this prospective question, we considered how reliability and validity measures would have changed if we had fewer trials per child. We filtered the data to include only children who had contributed at least N trials (for different values of N) and then repeatedly subsampled their trials and calculated our metrics on each subset. Datasets vary

in how many trials they have, so different minimum trials include different distributions and datasets, items, and children.

The results for accuracy and RT are shown in Figure 7. ICC reliability is not sensitive to the number of trials per subject, so it is not shown. For accuracy, more trials per administration increases test-retest reliability, up to 20 trials (the highest we measured), although the gains per additional trial decrease with more trials. Validity with CDI measures also increased as trials increased, with smaller gains past 5 trials. Based on accuracy, we recommend at least 10 trials per condition in general, and as many as possible for individual differences studies where test-retest reliability is essential.

Similar considerations apply for RT, but at lower total trial numbers. Nearly every trial a child does results in a usable accuracy measure, but only around 1/3 of trials result in a valid RT measure in our data. For test-retest reliability, having at least 5 RT measures per administration leads to notable improvements; gains beyond 5 are hard to measure because of the low number of subjects with sufficient trials, but more trials appears to be beneficial. For CDI measures, there are substantial gains for having 2 trials compared to 1 and continued benefits to including additional trials. Overall, for measuring individual differences in processing speed, designing studies to include as many trials as possible appears to be central to adequate measurement.

Recommendation 6: Match target and distractor images based on animacy.

One concern with LWL trial design is that children may have preferences for which image to look at separate from the auditory stimuli, and some of their looking behavior may be independent of the auditory stimulus. One approach is to decrease the preference discrepancy between a pair of images, so that in the absence of auditory cues, children prefer to look at both images approximately equally. While there are a number of factors that could play into children's image preferences (Fernald et al., 2008; Nothdurft, 2000; Pomper & Saffran, 2019), one factor that stood out in the data was animacy. Children have a strong preference for looking at animals and humans. When trials have one animate and one inanimate, children look more at the animate in the pre-stimulus period (Figure 8A) and are slower to respond when the inanimate is named. In contrast, for trials with target and distractor items matched on animacy, children's early looking patterns for animate and inanimate targets are highly similar. This difference is particularly relevant for the younger children, whose stimulus-contingent looking behavior is weakest, so more of the looking behavior is driven by animacy; however, effects of animacy were apparent in children up to the age of three (Figure 8B).

How much can this matter?

To quantify the potential cumulative consequences of different combinations of choices, we computed the difference in reliability when making all of the recommended choices with respect to data processing to an instance in which a researcher made a set of unfortunate choices that are nevertheless consistent with the literature. Consider two researchers analyzing accuracy in the same LWL dataset, Norm and Rex. Researcher Norm makes a set




Figure 8. Looking time profiles for different combinations of target and distractor animacy. Panel A shows overall; Panel B shows split by age category.




Figure 9. Differences in dependent measures between the recommended accuracy measurement and a sensible-sounding, but unfortunate alternative set of parameters.

of reasonable choices based on practices observed in other LWL papers, choosing to analyze a short analysis time window (300-1800 ms); apply baseline correction (baseline window: -2000-0 ms); exclude trials on which infants are not looking at the screen for at least 60% of the critical window or not looking at the screen during noun onset for insufficient attentiveness; and require that infants contribute at least 50% of all trials for inclusion. Researcher Rex, on the other hand, analyzes the same data and instead chooses parameters based on the data-driven recommendations above: a critical window of 600-4000ms, no baseline correction, and including all trials and sessions. Researcher Rex would gain substantial increases in ICC (0.2 [0.13, 0.27]) and test-retest reliability (0.21 [0.13, 0.29]).

Additionally, Rex could expect small increases in validity for correlations with comprehension (0.07 [-0.03, 0.17]) and production (0.07 [0, 0.13]). Despite both researchers making choices consistent with past practices, we would expect Rex to derive a measure with significantly improved psychometric properties compared to Norm, and therefore be in a substantially better position to study individual differences and longitudinal change.

General Discussion

Children’s early word recognition is a central construct for the study of language learning, so improving recognition measures can have widespread benefits for our understanding of language development. How can we maximize reliable, valid signal in the analysis of data in the LWL paradigm? We took a data-driven approach to this question, using a large-scale dataset of infant LWL data to quantify the impact of different analysis choices on measurement properties, leading to six core recommendations for data processing and experimental design.

For data processing, our strongest take-homes concern window choice and baseline correction. First, longer time windows (e.g., windows that end around 3000 ms after noun onset) maximize the reliability of accuracy measures, with no negative consequences for validity. While average accuracy tends to peak earlier in the timecourse of a trial (typically 1000-2000 ms after noun onset), these findings suggest that later looking windows capture signal that boosts the consistency of word recognition measures. Second, baseline correction, while an intuitive solution to imbalances in the saliency of images, causes large decrements in reliability and moderate decrements to validity, and we therefore recommend avoiding this practice. Finally, we formulate recommendations for choosing reaction time cutoffs (around 400ms tends to maximize reliability) and find that implementing trial- and session-level exclusions based on the amount of data contributed tend to have minimal consequences. Overall, choices about windows of analysis, approaches to exclusions, and how scores are computed can all have meaningful impacts on measure reliability.

One question that emerges from our analysis concerns the purity of the construct that LWL measures. In our analysis, long accuracy windows increased reliability more than validity. One speculative reason for this difference is that signal from later in the timecourse might be less coupled to language per se, and might instead reflect more general features of attention. Future research could use this difference as a starting point to further parcellate the attentional and linguistic contributors to language-driven eye-movements.

A more general take-home from our data processing analyses is that looking at patterns in grand-average time courses is not a trustworthy guide to child- or item-level psychometric properties. Past recommendations for time window selection have sometimes been motivated by inspecting consistent patterns of looking in grand-averages and noting when looking begins to decline across participants. These heuristics — while based on important patterns in looking behavior — may not always lead to optimal choices with respect to core measurement properties such as reliability. Wherever possible, measurement choices be guided by direct estimates of psychometric properties in order to maximize the

signal in infant looking behavior (Byers-Heinlein et al., 2022; Sander-Montant et al., 2025; Weaver et al., 2026).

Among our design-focused recommendations, the most striking result is the impact of animacy on infant looking behavior. When an image of an animate referent was paired with an image of an inanimate referent, infant looking was systematically biased towards the animate image. This finding has important consequences for estimating word recognition: accuracy for animate targets is substantially overestimated and accuracy for inanimate targets becomes substantially underestimated when animacy is not balanced in the target-distractor pairing. The results are especially dramatic for younger infants (under 18 months): looking to familiar inanimate targets can appear to be at chance when paired with an animate distractor. In other words, an improper balancing of saliency could make an infant appear to not know a word, even though they do. These results emphasize the importance of balancing saliency in image pairings as much as possible in LWL designs. This is especially true given that many of the current analytic approaches to correcting for saliency imbalances post hoc rely on difference-based measures (baseline correction; item-pair correction) which decrease reliability.

Limitations & Generalizability

The scope of our recommendations is fundamentally limited by the characteristics of the Peekbank dataset. First, while the Peekbank database is by far the largest dataset of its kind, it is not a comprehensive or representative set of LWL studies. To highlight just a few of these limitations, (1) the data is predominantly English word recognition collected from infants living in the U.S., and (2) the dataset is a convenience sample from labs who made their data openly available or otherwise accessible on request from the Peekbank team. LWL studies as a whole are likely more heterogeneous than what is represented in Peekbank — the datasets considered here come predominantly from a small group of labs or networks of researchers with similar training in LWL data collection.

At the same time, there is also substantial heterogeneity in the studies included in Peekbank, due to many studies pursuing different research questions and therefore making a variety of distinct design choices. To mitigate this concern, our analyses focused exclusively on standard trials that present familiar words with familiar images with limited or no targeted manipulations of the stimuli. The main benefit of focusing on standard trials is that it allows for a clear interpretation of reliability and validity — trials are truly designed to measure word recognition, rather than some other aspect of language processing. We view the remaining heterogeneity of study properties (e.g., specific words; differences in age) as a strength of our approach — our findings generalize broadly across a wide range of items and ages. However, by focusing on standard trials only, we also limit our conclusions to measuring familiar word recognition; our findings do not address how data processing and design practices would generalize to other estimation goals, e.g., estimating the magnitude of differences between two experimental conditions (though there is good reason to think that some of the more consequential choices, e.g. baseline correction, are likely have a similar impact).

While we summarize the choices that typically maximize reliability, we note that we often see a range of choices with similar outcomes in our analyses. For example, in determining the best time window for computing accuracy, a choice of ending a time window at 2800 vs. 3000 vs. 3200 ms leads to very similar outcomes. There is no one choice of time window that we can definitively infer as the “best” for maximizing reliability – instead, we see a pattern of graded improvements in reliability that follow a predictable pattern. For example, longer time windows tend to increase reliability holding the onset of the accuracy window constant. While we sometimes articulate which choices lead to the highest reliability descriptively, these are at best approximations. We offer these windows as data-driven default choices (especially in cases in which the decision would otherwise be ad hoc).

Conclusions

Experimental design and data processing decisions can have profound impacts on the properties of infant measures. Our findings show how compiling large-scale datasets such as Peekbank can help us estimate the impact of a wide variety of experimenter decisions, generalizable across a wide range of ages, stimuli, and experimental contexts. Establishing consistent practices and principles for reliable and valid measurement of infant looking sets the stage for asking central questions about how language skill develops.

References

- Adams, K. A., Marchman, V. A., Loi, E. C., Ashland, M. D., Fernald, A., & Feldman, H. M. (2018). Caregiver talk and medical risk as predictors of language outcomes in full term and preterm toddlers. *Child Development*, 89(5), 1674–1690.
- Bacon, D., & Saffran, J. (2022). Role of speaker gender in toddler lexical processing. *Infancy : The Official Journal of the International Society on Infant Studies*, 27(2), 291–300.
- Bergelson, E., & Aslin, R. (2017). Semantic Specificity in One-Year-Olds' Word Comprehension. *Language Learning and Development*, 13(4), 481–501. <https://doi.org/10.1080/15475441.2017.1324308>
- Bergelson, E., & Swingle, D. (2012). At 6-9 months, human infants know the meanings of many common nouns. *Proceedings of the National Academy of Sciences*, 109(9), 3253–3258.
- Bergmann, C., Tsuji, S., Piccinini, P. E., Lewis, M. L., Braginsky, M., Frank, M. C., & Cristia, A. (2018). Promoting replicability in developmental research through meta-analyses: Insights from language acquisition research. *Child Development*, 89(6), 1996–2009.
- Borovsky, A., & Peters, R. E. (2019). Vocabulary size and structure affects real-time lexical recognition in 18-month-olds. *PLOS ONE*, 14(7), e0219290. <https://doi.org/10.1371/journal.pone.0219290>
- Byers-Heinlein, K., Bergmann, C., & Savalei, V. (2022). Six solutions for more reliable infant research. *Infant and Child Development*, 31(5), e2296. <https://doi.org/10.1002/icd.2296>
- Csibra, G., Hernik, M., Mascaró, O., Tatone, D., & Lengyel, M. (2016). Statistical treatment of looking-time data. *Developmental Psychology*, 52(4), 521–536.
- Fernald, A., & Marchman, V. A. (2012). Individual differences in lexical processing at 18 months predict vocabulary growth in typically developing and late-talking toddlers. *Child Development*, 83(1), 203–222.
- Fernald, A., Perfors, A., & Marchman, V. A. (2006). Picking up speed in understanding: Speech processing efficiency and vocabulary growth across the 2nd year. *Developmental Psychology*, 42(1), 98–116. <https://doi.org/10.1037/0012-1649.42.1.98>
- Fernald, A., Pinto, J. P., Swingle, D., Weinberg, A., & McRoberts, G. W. (1998). Rapid gains in speed of verbal processing by infants in the 2nd year. *Psychological Science*, 9(3), 228–231.
- Fernald, A., Zangl, R., Portillo, A. L., & Marchman, V. A. (2008). Looking while listening: Using eye movements to monitor spoken language comprehension by infants and young children. In I. A. Sekerina, E. M. Fernández, & H. Clahsen (Eds.), *Language Acquisition and Language Disorders* (Vol. 44, pp. 97–135). Amsterdam: John Benjamins Publishing Company. <https://doi.org/10.1075/lald.44.06fer>
- Flake, J. K., & Fried, E. I. (2020). Measurement Schmeasurement: Questionable Measurement Practices and How to Avoid Them. *Advances in Methods and Practices in Psychological Science*, 3(4), 456–465. <https://doi.org/10.1177/2515245920952393>
- Frank, M. C., Braginsky, M., Yurovsky, D., & Marchman, V. A. (2017). Wordbank: An open repository for developmental vocabulary data. *Journal of Child Language*, 44(3), 677–694.
- Frank, M. C., Braginsky, M., Yurovsky, D., & Marchman, V. A. (2021). *Variability and*

- consistency in early language learning: *The wordbank project*. Cambridge, MA: MIT Press.
- Frank, M. C., Marchman, V. A., Bergey, C. A., Boyce, V., Braginsky, M., Kachergis, G., ... Zettersten, M. (2025). Continuous developmental changes in word recognition support language learning across early childhood. *eLife*. <https://doi.org/10.7554/eLife.109636.1>
- Garrison, H., Baudet, G., Breitfeld, E., Aberman, A., & Bergelson, E. (2020). Familiarity plays a small role in noun comprehension at 12–18 months. *Infancy : The Official Journal of the International Society on Infant Studies*, 25(4), 458–477.
- Golinkoff, R. M., Ma, W., Song, L., & Hirsh-Pasek, K. (2013). Twenty-five years using the intermodal preferential looking paradigm to study language acquisition: What have we learned? *Perspectives on Psychological Science*, 8(3), 316–339.
- Hirsh-Pasek, K., & Golinkoff, R. M. (1996). The intermodal preferential looking paradigm: A window onto emerging language comprehension. In *Language, Speech, and Communication. Methods for assessing children's syntax* (pp. 105–124). Cambridge, MA, US: The MIT Press.
- Hurtado, N., Marchman, V. A., & Fernald, A. (2007). Spoken word recognition by Latino children learning Spanish as their first language. *Journal of Child Language*, 34(2), 227–249. <https://doi.org/10.1017/S0305000906007896>
- Kominsky, J. F., Lucca, K., Thomas, A. J., Frank, M. C., & Hamlin, J. K. (2022). Simplicity and validity in infant research. *Cognitive Development*, 63, 101213. <https://doi.org/10.1016/j.cogdev.2022.101213>
- Law, F., II, & Edwards, J. R. (2015). Effects of Vocabulary Size on Online Lexical Processing by Preschoolers. *Language Learning and Development*, 11(4), 331–355. <https://doi.org/10.1080/15475441.2014.961066>
- Loi, E. C., Marchman, V. A., Fernald, A., & Feldman, H. M. (2017). Using eye movements to assess language comprehension in toddlers born preterm and full term. *The Journal of Pediatrics*, 180, 124–129. <https://doi.org/10.1016/j.jpeds.2016.10.004>
- López Pérez, M., Moore, C., Sander-Montant, A., & Byers-Heinlein, K. (2024). Infants' Knowledge of Individual Words: Investigating Links Between Parent Report and Looking Time. *Infancy*, 30(1), e12641. <https://doi.org/10.1111/inf.12641>
- Lord, F. M. (1956). The Measurement of Growth. *Educational and Psychological Measurement*, 16(4), 421–437. <https://doi.org/10.1177/001316445601600401>
- Mahr, T., McMillan, B. T. M., Saffran, J. R., Ellis Weismer, S., & Edwards, J. (2015). Anticipatory coarticulation facilitates word recognition in toddlers. *Cognition*, 142, 345–350. <https://doi.org/10.1016/j.cognition.2015.05.009>
- Marchman, V. A., Adams, K. A., Loi, E. C., Fernald, A., & Feldman, H. M. (2016). Early language processing efficiency predicts later receptive vocabulary outcomes in children born preterm. *Child Neuropsychology*, 22(6), 649–665. <https://doi.org/10.1080/09297049.2015.1038987>
- Marchman, V. A., Dale, P., & Fenson, L. (2023). *MacArthur-Bates communicative development inventories user's guide and technical manual*. Brookes Publishing.
- Marchman, V. A., Loi, E. C., Adams, K. A., Ashland, M., Fernald, A., & Feldman, H. M. (2018). Speed of language comprehension at 18 months old predicts school-relevant outcomes at 54 months old in children born preterm. *Journal of Developmental & Behavioral Pediatrics*, 39(3), 246–253.

- Mirman, D., Dixon, J. A., & Magnuson, J. S. (2008). Statistical and computational models of the visual world paradigm: Growth curves and individual differences. *Journal of Memory and Language*, 59(4), 475–494. <https://doi.org/10.1016/j.jml.2007.11.006>
- Moore, C., & Bergelson, E. (2024). Wordform variability in infants' language environment and its effects on early word learning. *Cognition*, 245, 105694. <https://doi.org/10.1016/j.cognition.2023.105694>
- Nothdurft, H.-C. (2000). Salience from feature contrast: Additivity across dimensions. *Vision Research*, 40(10-12), 1183–1201.
- Novack, M. A., Dworak, E. M., Han, Y. C., Kaat, A. J., Ustsinovich, V., Saffran, J., ... Gershon, R. C. (2025). Development and validation of the NIH Baby Toolbox® language measures. *Infant Behavior and Development*, 80, 102121. <https://doi.org/10.1016/j.infbeh.2025.102121>
- Olson, R. H., Pomper, R., Potter, C. E., Hay, J. F., Saffran, J. R., Ellis Weismer, S., & Lew-Williams, C. (2020). *Peyecoder: An open-source program for coding eye movements*. <https://doi.org/10.5281/zenodo.3939234>
- Peelle, J. E., & Van Engen, K. J. (2021). Time Stand Still: Effects of Temporal Window Selection on Eye Tracking Analysis. *Collabra: Psychology*, 7(1), 25961. <https://doi.org/10.1525/collabra.25961>
- Perry, L. K., & Saffran, J. R. (2017). Is a pink cow still a cow? Individual differences in toddlers' vocabulary knowledge and lexical representations. *Cognitive Science*, 41(4), 1090–1105. <https://doi.org/10.1111/cogs.12370>
- Peter, M. S., Durrant, S., Jessop, A., Bidgood, A., Pine, J. M., & Rowland, C. F. (2019). Does speed of processing or vocabulary size predict later language growth in toddlers? *Cognitive Psychology*, 115, 101238.
- Pomper, R., & Saffran, J. R. (2016). Roses are red, socks are blue: Switching dimensions disrupts young children's language comprehension. *PloS One*, 11(6), e0158459. <https://doi.org/10.1371/journal.pone.0158459>
- Pomper, R., & Saffran, J. R. (2019). Familiar object salience affects novel word learning. *Child Development*, 90(2), e246–e262. <https://doi.org/10.1111/cdev.13053>
- Potter, C. E., Fourakis, E., Morin-Lessard, E., Byers-Heinlein, K., & Lew-Williams, C. (2019). Bilingual toddlers' comprehension of mixed sentences is asymmetrical across their two languages. *Developmental Science*, 22(4), e12794. <https://doi.org/10.1111/desc.12794>
- Potter, C., & Lew-Williams, C. (unpublished). *Behold the canine!: How does toddlers' knowledge of typical frames and familiar words interact to influence their sentence processing?*
- Reznick, J. S. (1990). Visual preference as a test of infant word comprehension. *Applied Psycholinguistics*, 11(2), 145–166.
- Rhemtulla, M., & Bottesini, J. G. (2022). On the reliability of reliability estimates in infancy research. *Infant and Child Development*, 31(5), e2327. <https://doi.org/10.1002/icd.2327>
- Rogosa, D. R., & Willett, J. B. (1983). Demonstrating the Reliability the Difference Score in the Measurement of Change. *Journal of Educational Measurement*, 20(4), 335–343. <https://doi.org/10.1111/j.1745-3984.1983.tb00211.x>
- Ronfard, S., Wei, R., & Rowe, M. L. (2022). Exploring the linguistic, cognitive, and social skills underlying lexical processing efficiency as measured by the looking-while-listening

- paradigm. *Journal of Child Language*, 49(2), 302–325.
<https://doi.org/10.1017/S0305000921000106>
- Sander-Montant, A., Fibla, L., Byers-Heinlein, K., & Killam, H. (2025). *Keeping an eye on looking measures: Towards more robust developmental methods*. OSF.
https://doi.org/10.31234/osf.io/vxu53_v1
- Sander-Montant, A., Pérez, M. L., & Byers-Heinlein, K. (2023). The more they hear the more they learn? Using data from bilinguals to test models of early lexical development. *Cognition*, 238, 105525. <https://doi.org/10.1016/j.cognition.2023.105525>
- Schreiner, M. S., Zettersten, M., Bergmann, C., Frank, M. C., Fritzsche, T., Gonzalez-Gomez, N., et al.others. (2024). Limited evidence of test-retest reliability in infant-directed speech preference in a large preregistered infant experiment. *Developmental Science*, 27(6), e13551. <https://doi.org/10.1111/desc.13551>
- Snedeker, J. (2013). Children’s sentence processing. In *Sentence Processing*. Psychology Press.
- Swingley, D., & Aslin, R. N. (2000). Spoken word recognition and lexical representation in very young children. *Cognition*, 76(2), 147–166.
[https://doi.org/10.1016/s0010-0277\(00\)00081-0](https://doi.org/10.1016/s0010-0277(00)00081-0)
- Swingley, D., & Aslin, R. N. (2002). Lexical Neighborhoods and the Word-Form Representations of 14-Month-Olds. *Psychological Science*, 13(5), 480–484.
<https://doi.org/10.1111/1467-9280.00485>
- Trafimow, D. (2015). A defense against the alleged unreliability of difference scores. *Cogent Mathematics*, 2(1), 1064626. <https://doi.org/10.1080/23311835.2015.1064626>
- Venker, C. E., Pomper, R., Mahr, T., Edwards, J., Saffran, J., & Ellis Weismer, S. (2020). Comparing automatic eye tracking and manual gaze coding methods in young children with autism spectrum disorder. *Autism Research*, 13(2), 271–283.
<https://doi.org/https://doi.org/10.1002/aur.2225>
- Von Holzen, K., & Mani, N. (2012). Language nonselective lexical access in bilingual toddlers. *Journal of Experimental Child Psychology*, 113(4), 569–586.
<https://doi.org/10.1016/j.jecp.2012.08.001>
- Wass, S. V., Smith, T. J., & Johnson, M. H. (2013). Parsing eye-tracking data of variable quality to provide accurate fixation duration estimates in infants and adults. *Behavior Research Methods*, 45, 229–250.
- Weaver, H., & Saffran, J. R. (2026). Interrogating Early Word Knowledge: Factors That Influence the Alignment Between Caregiver-Report and Experimental Measures. *Developmental Science*, 29(1), e70088. <https://doi.org/10.1111/desc.70088>
- Weaver, H., Saffran, J., & Arunachalam, S. (2026). Looking forward: Eye-gaze methods in vocabulary development research. *Trends in Cognitive Sciences*.
- Weaver, H., Zettersten, M., & Saffran, J. R. (2024). Becoming word meaning experts: Infants’ processing of familiar words in the context of typical and atypical exemplars. *Child Development*, 95(5), e352–e372. <https://doi.org/10.1111/cdev.14120>
- Weisleder, A., & Fernald, A. (2013). Talking to children matters: Early language experience strengthens processing and builds vocabulary. *Psychological Science*, 24(11), 2143–2152.
<https://doi.org/10.1177/0956797613488145>
- Yurovsky, D., Wade, A., Kraus, A., Gengoux, G., Antonio, Y., & Frank, M. (2017). *Developmental changes in the speed of social attention in early word learning*.

unpublished manuscript.

Zettersten, M., Yurovsky, D., Xu, T. L., Uner, S., Tsui, A. S. M., Schneider, R. M., . . .

Frank, M. C. (2022). Peekbank: An open, large-scale repository for developmental eye-tracking data of children's word recognition. *Behavior Research Methods*, 55(5),

2485–2500. <https://doi.org/10.3758/s13428-022-01906-4>