

1 Data-driven recommendations for increasing reliability in measures of infant word recognition

2 Veronica Boyce^a, Virginia A. Marchman^a, Claire Augusta Bergey^b, Mika Braginsky^a, George
3 Kachergis^a, Federico Giovannetti^c, Jess Mankewitz^d, Stephan Meylan^e, Ben Prystawski^a,
4 Robert Z. Sparks^a, Adrian Steffan^f, Alvin Wei Ming Tan^a, Tian Linger Xu^g, Michael C.
5 Frank^a, & Martin Zettersten^h

6 ^a Department of Psychology, Stanford University, Stanford, CA, USA

7 ^b Department of Linguistics, Stanford University, Stanford, CA, USA

8 ^c Unidad de Neurobiología Aplicada (UNA, CEMIC-CONICET), Buenos Aires, Argentina

9 ^d Department of Psychology, University of Wisconsin, Madison, WI, USA

10 ^e Department of Linguistics, University of California, Berkeley, CA, USA

11 ^f Department of Psychology, Ludwig Maximilian University, Munich, Germany

12 ^g Department of Psychology, Indiana University, Bloomington, IA, USA

13 ^h Department of Cognitive Science, University of California, San Diego, CA, USA

14 Author Note

15 TODO UPDATE AUTHOR LIST AND AFFILIATIONS

Abstract

Children’s early word recognition is both an important phenomenon in its own right and a key index of early language more broadly. The looking-while-listening paradigm (LWL), in which eye-movements to the referent of a word are tracked to reveal online language comprehension, is a key paradigm for measuring word recognition in young children. Yet the interpretation and use of this paradigm has been hindered by variation in how LWL data are analyzed. Here we use data from Peekbank, a repository of LWL data from studies with almost 4000 children (6 months – 5 years), to search for analytic and design decisions that maximize the reliability and validity of LWL measures. Based on these analyses, we recommend the use of long analytic windows for computing accuracy, and we do not find benefits for correcting for baseline preferences or excluding data from children or trials with low data. In our dataset, this combination of analytic choices increases reliability by approximately .20 and validity by .07. We also recommend that LWL paradigms be designed with a minimum of 10 target trials per child per condition and that experimentalists match their stimuli for animacy. We hope that these recommendations lead to improved and unified standards for the measurement of children’s word recognition.

Keywords: keywords

Word count: X

Data-driven recommendations for increasing reliability in measures of infant word recognition

Introduction

Over the first few months of life, children become experts in their native language, organizing the phonological system of the speech that they hear and parsing speech into meaningful units. A bit later, children begin to show early signs of comprehension of words or phrases through simple behaviors such as turning their head when their name is called (e.g., “Hey, Billy!”) or reaching for an object that is named (e.g., “Where’s the ball?”). These everyday accomplishments mark important milestones in children’s development. Yet such global assessments may miss continuous processes of growth in underlying skills and the millisecond-level processes that take place during real-time language comprehension. Children, like adults, process speech incrementally, as it unfolds in time (Fernald, Pinto, Swingley, Weinberg, & McRoberts, 1998; Frank et al., 2025; Snedeker, 2013), and use features of the speech signal to extract meaning that will support language comprehension (Law & Edwards, 2015; Weisleder & Fernald, 2013).

These moment-to-moment processes are difficult to document without child-friendly procedures that allow precision in capturing how children’s behaviors are time-locked to speech. In particular, the “looking while listening” (LWL) procedure has emerged as a *de facto* standard for measuring language in infants and children, especially when they are younger than two years old, as other options for direct assessment during this period are limited (Frank, Braginsky, Yurovsky, & Marchman, 2021).¹ In the standard version of LWL, children see two pictures on a screen and hear a referring expression, e.g. “look at the ball”; their time-locked eye-movements to the ball over the distractor image are then taken to be evidence for language comprehension (Fernald et al., 2008).

LWL has been a key tool for enhancing our understanding of infants’ language processing and knowledge (Golinkoff, Ma, Song, & Hirsh-Pasek, 2013). For example, using LWL methods, we can measure increases in speed and accuracy in responding to familiar words (Fernald, Perfors, & Marchman, 2006; Frank et al., 2025) and increases in the detail of their phonological representations (Moore & Bergelson, 2024; Swingley & Aslin, 2000) during the first years after birth. Moreover, measures of children’s skill in the LWL task have also proven useful as an index of individual differences. For example, children who show more efficient language processing at 18 months show larger vocabulary size and more rapid vocabulary growth over the second and third years than children with less efficient processing (Fernald & Marchman, 2012; Fernald et al., 2006; Frank et al., 2025; Peter et al., 2019). Language processing efficiency has also been shown to link to later language, cognitive, and non-verbal outcomes in both typically-developing (Law & Edwards, 2015; Peter et al., 2019) and preterm children (Virginia A. Marchman, Adams, Loi, Fernald, & Feldman, 2016; V. A. Marchman, Dale, & Fenson, 2023; Virginia A. Marchman et al., 2018).

¹ Here we adopt LWL as a convenient label while recognizing that the procedure goes by a number of different names in the literature including the “intermodal preferential looking procedure” (Fernald, Zangl, Portillo, & Marchman, 2008; Hirsh-Pasek & Golinkoff, 1996; Reznick, 1990).

One key reason for the success of this research is the care that has been invested in the development of the looking-while-listening method. Past work has articulated best practices in collecting and analyzing looking-while-listening measures, with a particular focus on stimulus and design choices that make looking preferences maximally informative about word recognition processes (Fernald et al., 2008). However, despite widespread use of the looking-while-listening method in language development research, there is still limited consensus on how to best process and analyze time-series data of this type [from either adults or infants; Mirman, Dixon, and Magnuson (2008); Bergelson and Swingley (2012); Peelle and Van Engen (2021); Von Holzen and Mani (2012); Weaver, Saffran, and Arunachalam (2026)].

In the developmental field especially, this lack of consensus likely stems from structural factors: infant recruitment is time- and cost-intensive and datasets collected by individual labs are often small (Bergmann et al., 2018; Frank, Braginsky, Yurovsky, & Marchman, 2017). These issues make it difficult for any one researcher to systematically evaluate the consequences of specific analytic choices independently of the data collected for a particular study (cf. Peelle & Van Engen, 2021). Instead, researchers are often forced to make ad-hoc decisions or to adopt idiosyncratic lab practices that themselves originate from subjective judgments of what “seems to work”. By decreasing uniformity in the literature, this variation also increases the challenge of triangulating what precisely is being measured by LWL (Weaver et al., 2026). How can we put the measurement and analysis of infant word recognition and online language processing on firmer empirical footing?

Peekbank is a new data resource that offers a potential solution to this core issue by providing a large database of infant eye-tracking data from the looking-while-listening paradigm for exploring potential methodological questions (Zettersten et al., 2022). The most recent release of Peekbank at time of writing (version 2026.1), currently houses 44 datasets from infants across a wide age range [6 months - 5 years] and across a broad range of test words and materials, resulting in observations from over 3915 infants participating in the looking-while-listening task. The goal of the current manuscript is to use these data to systematically explore key analytic decisions that researchers make when using this paradigm, offering best-practices guidance for measurement.

Our approach throughout is to seek analytic choices that distill the best possible measures from the rich time course of behavior in the LWL task. Weaver et al. (2026) highlight open questions about both the reliability and validity of LWL. Starting with reliability, we aim for analyses that yield the maximal signal about individual children relative to both measurement error and irrelevant artifacts (Byers-Heinlein, Bergmann, & Savalei, 2022; Rhemtulla & Bottesini, 2022; Schreiner et al., 2024). Here, to measure reliability, we use both within-dataset inter-item reliability (ICC) and test-retest reliability within datasets with multiple LWL sessions close in time for the same children.

Turning to validity, we want to make analytic decisions that maximize the relationship with other measures that are thought to reflect related constructs (Flake & Fried, 2020; Kominsky, Lucca, Thomas, Frank, & Hamlin, 2022). Here, we select as our primary measure the correlations between LWL measures and parent reports of vocabulary measured with the

MacArthur-Bates Communicative Development Inventories [MB-CDI, a popular parent report measure of early vocabulary that is widely used in conjunction with LWL tasks; Frank et al. (2021); Fenson et al. (1994)]. We focus specifically on the relationships between aggregate measures of vocabulary size from the MB-CDI and LWL data, given both the documented relationships between these (reviewed in V. A. Marchman et al., 2023) and the mixed evidence on item-level alignment between LWL and MB-CDI data (López Pérez, Moore, Sander-Montant, & Byers-Heinlein, 2024; Weaver & Saffran, 2026). This correspondence reflects the general relationship between word recognition in the moment (indexed by LWL) and the growth of early vocabulary over time (indexed by the MB-CDI) (Frank et al., 2025).

Our focus in this paper is on providing guidance for navigating some of the many decisions facing an investigator analyzing LWL data. Recent work using an earlier version of Peekbank (as well as other datasets) has investigated the psychometric properties of a range of different dependent measures in LWL tasks and their general construct structure (Sander-Montant, Fibla, Byers-Heinlein, & Killam, 2025). We focus instead on the consequences of varying analytic choices for two main measures, which tend to have the strongest psychometric properties (Sander-Montant et al., 2025). The first is *accuracy*: how much of the time a child looks at the target image compared to looking at either the target or distractor image over some period of time after the auditory stimulus. Children’s responses unfold over time, but it is not known what start and end time maximizes the signal of children’s response to the stimulus (Peelle & Van Engen, 2021). The second measure we focus on is *reaction time* – how long it takes a child to change where they are looking in response to the stimulus (Fernald et al., 2008). There are even more analytic options available for this metric – whether to measure to a child leaving the distractor or looking at the target, whether to measure RT on a log or raw scale, and how long after the start of the stimulus to start counting the RTs. For both of these measures, there are also questions about whether and when to exclude either children or trials.

While the LWL task is used for a wide variety of purposes, including measurements of individual differences (e.g., Peter et al., 2019) as well as measurements of experimentally-induced condition differences (e.g., Swingley & Aslin, 2000), here we focus only on what we refer to as “vanilla” word recognition items. These are items that test vocabulary directly with familiar targets, familiar distractors, and no informative carrier phrases or experimental manipulations. We focus on these items both because they are the most commonly used LWL items and because they are the case in which reliability and validity is most clearly defined. Vanilla trials are expected to reflect the construct of word knowledge and correlate with other measures of word knowledge, including vocabulary size.

The plan of the paper is as follows. We begin by overviewing the subset of the Peekbank database we selected for our analyses. We then make a set of targeted recommendations, providing evidence for each based on the impact of our recommendation on the average reliability and validity of the resulting measures. Our first set of recommendations focus on data processing decisions and our second set of recommendations concern design considerations. To foreshadow our results, we find that while many common

data analysis practices provide relatively reasonable levels of reliability and validity, it may be possible to further optimize reliability and validity through careful, data-driven decision-making.

Methods: Dataset and approach

Peekbank is an aggregation of datasets from a number of LWL studies. This work is based on Peekbank version 2026.1 (released February 2026) which has 134929 trials from 3915 children and 44 datasets.

Inclusion criteria

We focus our exploration of data processing choices on a set of items where we have the clearest expectations around what we are trying to measure. While LWL is also used to study a variety of experimental effects, such as the effects of noise, mispronunciations, novel word learning, and expectations of mutual exclusivity, these items may index different constructs. We limit our analyses to what we term “vanilla” trials: trials that represent standard LWL trials testing familiar words without additional experimental manipulations. For a trial to be considered “vanilla”, both the target and the distractor images must be of familiar objects. For the auditory stimulus, the target word must be a real word that occurs once and is the first and only source of information about the target. The carrier phrase must be a well-formed, grammatical phrase, without any switches in language, and without any intentionally added noise or audio filtering. Vanilla trials could come from studies with only vanilla trials, from the filler or baseline trials of experimental studies, or even from studies where the experimental conditions met the vanilla criteria. (See SI: Criteria for Inclusion for more detailed criteria used to handle borderline situations.) While still heterogeneous, these “vanilla” trials have a shared interpretation for how looking behavior is expected to relate to underlying language knowledge and processing. We call the start of the target word in the audio the stimulus onset and label this as the zero point for timing. We only include trials from children 60 months or younger (see Figure S1 for age distributions). We do not place restrictions on the language of the trials or children’s language background.

Descriptives

This paper is based on 36 datasets, comprising 3171 children, 4851 administrations, and 75636 trials. There are MB-CDI comprehension scores for 632 children from 12 of the datasets, and MB-CDI production scores for 1143 children from 20 of the datasets. Our test-retest data comes from 5 datasets with multiple administrations within 1.5 months of each other, covering 365 children (788 administration pairs). See Table 1 for a list of included datasets. Figure S2 has time course plots showing the number of trials and timing of available data for each dataset.

Approach

Throughout we use two reliability metrics and two validity metrics. Reliability refers to how consistently measurements are measuring the same thing. One of our measures of reliability is an intra-class correlation (ICC). Specifically, we use a measure of inter-rater

Table 1
Descriptives of the included datasets.

	Dataset Name	Pct Trials	Trials	Admins	Subjects	Age	CDI	Test-retest	Language
1	Adams et al. (2018)	18.4%	13899	711	69	22 [13 - 38]	x	x	English
2	Fernald & Marchman (2012)	15.6%	11828	679	122	22 [17 - 32]	x	x	English
3	Weaver et al. (2024)	6.5%	4948	248	141	16 [14 - 24]		x	English
4	Fernald et al. (2013)	5.8%	4359	178	80	20 [17 - 26]	x	x	English
5	Hurtado et al. (2008)	5.3%	4035	136	76	21 [17 - 27]	x		Spanish
6	Fernald et al. (2006)	4.9%	3692	229	63	20 [15 - 25]	x		English
7	Yurovsky et al. (2013)	3.7%	2810	413	413	34 [12 - 60]			English
8	Yurovsky et al. (2017)	3.1%	2350	327	327	37 [8 - 60]			English
9	Bergelson & Swingley (2012)	2.8%	2081	84	84	12 [6 - 21]	x		English
10	Kartushina & Mayor (2019)	2.7%	2072	69	69	7 [6 - 9]	x		Norwegian
11	Weaver & Saffran (2026)	2.4%	1788	69	69	19 [18 - 20]	x		English
12	Weisleder & Fernald (2013)	2.3%	1769	58	29	22 [18 - 27]			Spanish
13	Yurovsky & Frank (2017)	2.3%	1747	285	285	26 [13 - 59]			English
14	Borovsky & Peters (2019)	2.2%	1676	79	79	18 [17 - 20]	x		English
15	Potter & Lew Williams (2023)	2.1%	1626	67	67	24 [21 - 27]	x		English
16	Yoon et al. (2015)	2.0%	1500	199	199	43 [13 - 60]			English
17	Sander-Montant et al. (2023)	1.8%	1381	123	123	22 [12 - 31]	x		English, French
18	Mahr et al. (2015)	1.6%	1232	29	29	21 [18 - 24]	x		English
19	Egger et al. (2020)	1.4%	1052	54	54	18 [18 - 19]	x		Dutch
20	Garrison et al. (2020)	1.4%	1046	35	35	15 [12 - 18]	x		English
21	Hurtado et al. (2007)	1.3%	1010	49	49	24 [15 - 37]			Spanish
22	Pomper & Saffran (2017)	1.0%	785	82	82	17 [14 - 19]			English
23	Frank et al. (2016)	1.0%	781	108	108	35 [12 - 60]			English
24	Bacon & Saffran (2022)	1.0%	754	38	38	23 [22 - 24]	x		English
25	Ronfard et al. (2022)	1.0%	737	40	40	20 [18 - 24]	x		English
26	Perry & Saffan (2017)	0.9%	718	42	42	20 [19 - 22]	x		English
27	Casillas et al. (2017)	0.9%	651	23	23	32 [9 - 48]			Tseltal
28	Swingley & Aslin (2002)	0.8%	596	50	50	15 [14 - 16]	x		English
29	Pomper & Saffran (2016)	0.6%	471	60	60	44 [41 - 47]			English
30	Kremin et al. (2023)	0.5%	410	60	60	42 [37 - 49]			English, French, Spanish
31	Pomper & Saffran (2015)	0.5%	389	25	25	40 [38 - 42]			English
32	Moore & Bergelson (2022)	0.5%	384	29	29	18 [16 - 20]	x		English
33	Potter et al. (2019)	0.4%	309	44	23	23 [18 - 29]	x	x	English
34	Byers-Heinlein et al. (2017)	0.4%	284	48	48	20 [19 - 21]	x		English, French
35	Pomper & Saffran (2019)	0.3%	249	44	44	40 [38 - 43]			English
36	Pomper & Saffran (2018)	0.3%	217	37	37	39 [38 - 43]			English
Total		100%	75636	4851	3171	22 [6 - 60]	21	5	

reliability that measures (within a dataset) how much different items (target labels) are consistent with each other in “rating” the subjects (kid-administrations). See SI: Details about ICC for more specifics. Another measure of reliability is test-retest reliability, which measures the correlation for measures on the same child at two different administrations within 1.5 months of each other. This measure is only possible on a subset of the datasets that closely spaced repeat administrations. While we are limited in how much data we have for test-retest reliability, this is the metric that is most important for measuring stable individual differences.

In addition to having reliable measures, we also care that our measurements correspond to the construct of interest (vocabulary and language processing speed). For validity, we correlate with CDI production measures and CDI comprehension measures, from the datasets where we have CDI administrations. We use both short and long form CDIs, using the proportion of items a child was marked as knowing as our measure. In general, correlating with production and comprehension in the expected directions is desirable, although we also may expect that LWL measures are detecting distinct aspects of language learning compared to what is represented in the CDI (see Frank et al., 2025 for a detailed analysis).

All the code and data to reproduce the analyses is available at

<https://github.com/peekbank/peekbank-method> [TODO: anonymize for review if needed].

Data processing recommendations

Our first set of recommendations focuses on data processing decisions. Looking while listening data undergoes significant post-processing to go from the experiment to the data that is analysed. Experiment design, eye-tracking, and eye-gaze coding considerations are covered elsewhere (Fernald et al., 2008; Olson et al., n.d.; Venker et al., 2020; Wass, Smith, & Johnson, 2013), so we focus here on the process of going from a series of timepoints to the measures of accuracy and reaction time that are analysed. There are a number of decisions involved, including what windows and time cutoffs to use, and what exclusions to make around missing or low amounts of data.

Recommendation 1: Measure accuracy over a long window.

In LWL studies, accuracy is typically calculated as the fraction of time a child is looking at the target divided by the time they are looking at the target or the distractor. These fractions are computed over some window of time, but the selection of this window is complex. Typically windows start shortly after the point of disambiguation and continue until several seconds later; this window is intended to capture the peak of children’s stimulus-driven looks to the target. However, in practice studies vary in what windows are used, and often window selection is a locus for analytic flexibility (choosing windows post hoc to maximize a particular effect) (Pelle & Van Engen, 2021). Common analysis windows include starting 300 milliseconds after stimulus onset and ending 1800 ms after stimulus onset (following Fernald et al., 2008) and 367ms to 2000ms after stimulus onset (following Swingley & Aslin, 2002), although windows that extend until 3 or 4 seconds after stimulus onset are also used (Bergelson & Swingley, 2012; López Pérez et al., 2024; bergelson2017?). In this analysis, we begin from first principles and ask what windows maximize reliability and validity on word recognition trials.

Analytic approach. We first calculated the ICC reliability for accuracy across datasets, computing average ICC across a range of different window onsets and offsets. Next, we added the additional metrics of CDI correlations and test-retest reliability, which were calculated by dataset, for each combination of parameter values. To aggregate across datasets, we used a linear mixed effect model with random intercept by dataset, weighted by the number of administrations. We interpret the fixed-effect contrasts as what differences in outcome measure would be expected from switching from one set of parameters to another on a generic dataset. We report estimates and 95% confidence intervals.

The best data-processing options might be different for children of different ages, so we ran the same linear mixed effect models separately for data just on children in each of 4 age groups: <18 months, 18-24 months, 24-36 months, and >36 months.

Results. Average ICCs across windows are shown in Figure 1. As expected, including time points from before the point of disambiguation decreases reliability. Our initial exploration found a region of high reliability with onsets in the 600-800 range and offsets in the 2800-4000 range, surrounded by a gradual decline in ICC reliability.




Figure 1. A heat map showing the average ICC across datasets for different start and end points for the accuracy window. A region with the highest ICC (best reliability) is outlined in black, and the region most commonly represented in the literature is outlined in blue.




Figure 2. Per dataset outcome measures for two window options: 300ms-1800ms (red) and 600ms-4000ms (blue). Dot size is proportional within each panel to the number of administrations.

250 A dataset by dataset comparison of the differences between a 300-1800ms window
 251 (representative of common practice) and a 600-4000 ms window are shown in Figure 2.
 252 Nearly all datasets have higher ICC reliability for the later and longer window than the
 253 shorter window. All 5 datasets with close in time pairs of administrations also show
 254 increased reliability on this later and longer window. Correlations with CDI vary across

datasets, with an overall similar effect for both windows.

Our mixed effect model confirms that there are expected reliability gains from switching from a 300-1800ms window to a 600-4000ms one. ICC reliability is expected to increase by 0.12 [0.08, 0.16] and test-retest reliability by 0.07 [-0.02, 0.16]. There are limited differences in validity between these two windows (-0.02 [-0.07, 0.03] for comprehension, and 0.03 [-0.02, 0.08] for production).

There are a wide range of window options with generally comparable reliability and validity measures (Figure S4). For instance, a 400 - 3000 ms window is not significantly different from a 600 - 4000 ms window on any of the outcome measures.

These patterns generally hold across age. The overall level of reliability rises with age, as younger children have less signal, and experiments with younger children may also use fewer trials (Figure S3). However, to make data-driven recommendations, a question is whether there are substantive differences in what windows maximize validity and reliability across age.

While the strength of evidence varied, the overall pattern of results was remarkable consistent across the different age groups (Figure S5). For instance, across different ages there is an improvement in ICC for using a 600-4000 window compared to a 300-1800ms one. For <18 months, the ICC difference is 0.06 [-0.03, 0.16]; for 18-24 months, 0.13 [0.07, 0.19]; for 24-36 months, 0.06 [0.01, 0.12]; and for 36+ months 0.15 [0.06, 0.23]. The main difference that is seen across ages is that when there is less signal to begin with, there's less difference between different measurements.

For CDI production, there might be slight age effects, but the differences are not significant and are difficult to interpret. For <18 months, a longer window results in change of 0.02 [-0.07, 0.12] in correlation with CDI production ; for 18-24 months, -0.03 [-0.08, 0.02]; for 24-36 months, 0.06 [-0.03, 0.15]. These small and inconsistent age effects could be due to dataset differences across different ages or small but true developmental differences.

Exactly when the accuracy window starts is not important, although we recommend starting between 400ms and 600 milliseconds to optimally balance validity and reliability. Continuing for a long window also increases reliability.

Strong effects are likely to show up in any reasonable window, but for eking the most reliability out of data, we recommend windows that start somewhere in the 400-600 ms range and go for longer than 2000 ms. Why would later starting and later ending windows increase reliability? This suggests that there is some signal in the post 2000ms window, and that there may not be so much signal in the early 300-600ms time. While some children may look quickly to the target and then look away, other children who look to the target later or who continue looking at the target at still contributing signal in this region, and overall children in many datasets are in aggregate looking to the target more than the distractor more in this region (even if less than at the peak, Figure S2). Importantly, since accuracy is measured as the ratio of target looks to either target or distractor looks, looking completely away does

not count against accuracy.

Windowing seems to primarily affect reliability measures and not validity measures. It is not currently knowable whether the consistent later signal reflects child-level differences that are not language-relevant or whether these child-level differences are linguistic, but are distinct signal from the CDI.

Recommendation 2: Do not use baseline-corrected accuracy.

One commonly acknowledged issue in LWL studies is that looking behavior could be partly driven by preferences over the target images themselves (Fernald et al., 2008; Novack et al., 2025): Children may have a baseline preference to look at one image over another, even in the absence of auditory stimuli. Although thoughtful design choices can minimize these issues (see Recommendation 6), some researchers also attempt to deal with baseline differences in stimulus salience in their dependent measures. This practice is known as “baseline correction”: at the data-processing stage, researchers use a difference measure between accuracy on a post-stimulus window and “accuracy” (looks to target versus looks to target or distractor) on a pre-stimulus window (Bergelson & Swingley, 2012; Weaver, Zettersten, & Saffran, 2024). The pre-stimulus window can be the 2 seconds prior to disambiguation (Weaver et al., 2024) or the full 3-4 second pre-stimulus period (Bergelson & Swingley, 2012).

Analytic approach. We considered baseline-corrected accuracies with various pre-stimulus windows, focusing on a 600ms-4000ms main window (although the results generalize to other main windows). We restricted ourselves to datasets with at least some data from 2000ms or more before stimulus onset (see Figure S2 for which datasets have data when). We also included a “no baseline” version which was just accuracy on 600-4000ms with no correction. Baseline windows started at -1000, -2000, -3000 and -4000ms and ended at stimulus onset.

Results. Baseline-corrected accuracy led to worse reliability and validity than accuracy.

As an example, Figure 3 shows per-dataset results for accuracy compared to a baseline-corrected accuracy using a -2000 - 0 baseline window. All baseline correction options that we considered led to significantly worse reliability and validity measures compared to using normal accuracy without any correction. For a regular window of 600-4000, baseline correction is expected to make ICC worse by at least 0.1 [0.04, 0.16]; test-retest worse by at least 0.19 [0.08, 0.3]; correlation with comprehension worse by at least 0.04 [-0.05, 0.12] and correlation with production worse by at least 0.09 [0.05, 0.13].

Similar patterns hold for different accuracy windows (Figure S6) and across different age groups (Figure S8).

Why would incorporating baseline information lead to worse metrics? There is a complex literature on the psychometrics of difference scores and when they result in lowered reliability (e.g., Lord, 1956; Rogosa & Willett, 1983; Trafimow, 2015). The situation for looking while listening data is that children’s looking behavior is noisy but the signal across




Figure 3. Comparison between a baseline-corrected accuracy in red and accuracy in blue across datasets and measures. Both use an accuracy window of 600 - 4000ms, the baseline window was -2000 - 0ms.

baseline and critical portions of a trial is likely correlated (in part due to the very salience that the baseline correction is trying to remove). Unfortunately, it is under precisely these conditions – low reliability and high correlation between components – that difference scores exhibit the lowest reliability (Trafimow, 2015). Thus, while baseline-correction may lead to less biased grand average accuracy scores (by removing item salience), it also leads to less individual-level measures.

Some researchers address stimulus saliency by comparing across paired trials or adding additional predictors to their models, but these are out of scope; these are out of scope for this paper, but may be more appropriate. [citations; not sure what to say, but we should acknowledge somewhere]

Recommendation 3: Measure reaction times with a minimum reaction time of around 400 ms.

In addition to accuracy measures, LWL studies also provide the potential for computing children’s reaction time [RT; Fernald et al. (2008)]. RT is typically computed on trials when the child is initially looking at the distractor and shifts to the target after stimulus onset, as there is not a way to measure how quickly a child decides to keep looking at the target. RT computations include many choice points:

- **Launch vs. landing:** RTs can be computed based on when a child’s eyes leave from the distractor or when a child’s eyes land on the target.
- **Exclusion of short and long RTs:** Short and long RTs are often excluded because they are considered unlikely to be generated in response to the stimulus.

- **Linear vs. logarithmic analysis:** Reaction time data are often analyzed using a logarithmic transform (consistent with the idea that reaction times are log-normally distributed and that effects on them may be multiplicative rather than additive) (Csibra, Hernik, Mascaro, Tatone, & Lengyel, 2016; Frank et al., 2025).




Figure 4. Heat map of different parameter options for measuring RT and how they on average affected ICC.

Analytic approach. We considered a wide range of options for measuring RT, fulling crossing a) launch or land based, b) minimum RTs between 0 and 1000 ms, c) maximum RTs between 1600 and 4000, and d) raw or log scale.




Figure 5. Comparison between two different sets of RT parameters. In red, an example of common practice, using a launch-based RT from 300-1800ms. In blue, an RT that maximizes ICC, with a land RT from 400 to 4000ms, on the log scale.

Results. An initial exploration of ICC reliability (Figure 4) indicated that ICC reliability was maximized for minimum RTs around 400 ms, with longer max RTs, and with land based RT.

A comparison between the a typical way of calculating RT and an alternative that maximizes ICC are shown in Figure 5.

Compared to the “standard” launch-based RT restricted to 300-1800ms, a land-based minimum of 400 -4000 RT (logged) is expected to increase ICC by 0.11 [-0.01, 0.23] with little change in test-retest reliability (0.04 [-0.18, 0.27]). Differences in CDI correlation are not substantive (0.04 [-0.08, 0.17] for comprehension, and 0.03 [-0.06, 0.12] for production).

These patterns generally hold for children of different ages (Figure S9 and Figure S11).

Overall RT measures complement accuracy measures and capture substantial signal. A wide range of parameter options have similar properties (Figure S10 and Figure S12), but a land-based RT on 400 - 4000ms maximizes reliability without negatively impacting validity. Log-scale transformations should be considered, but may depend on modelling and theoretical considerations.

Recommendation 4: Trial-level and child-level exclusions for low amounts of data are generally unnecessary.

[why people exclude, range of practice] A common issue with developmental research is when to exclude trials for behavioral patterns that don’t seem to match the task and when to exclude children for contributing too few valid trials.

For LWL studies, trial-level exclusions tend to be either for contributing too few time points (ex. too much tracker loss or a child looking away from the screen too much). Some papers use heuristic exclusions for inattention prior to coding looking behavior (Adams et al., 2018), while others exclude trials that have too much looking away from the target and distractor. Cutoffs vary, with different experiments requiring minimum looking at the images at least 20%, 50% or 66% of the time (Borovsky & Peters, 2019; Pomper & Saffran, 2019; Weaver et al., 2024). Studies also occasionally filter for looking at the target or distractor specifically at the point of disambiguation (C. E. Potter, Fourakis, Morin-Lessard, Byers-Heinlein, & Lew-Williams, 2019).

If trials with fewer valid looking points still provide signal of the child-level language behavior, excluding them will decrease reliability and validity. On the other side, if these trials are not representing the construct and are just noise, excluding them will increase reliability and validity.

In addition to trial level exclusions, studies also do child-level exclusions based on the minimum number of valid trials. For accuracy measures, children are often excluded if they had fewer than 25% or 50% of the total trials, either within a condition, or overall (Adams et al., 2018; Garrison, Baudet, Breitfeld, Aberman, & Bergelson, 2020; Perry & Saffran, 2017).

For RT, it is instead standard to require a certain number of trials, often 2 (Adams et al., 2018; Fernald et al., 2006).

Analytic approach. We consider 3 types of exclusions: trial-level exclusions for accuracy, administration-level exclusions for accuracy, and administration-level exclusions for RT. We first explored trial-level accuracy exclusions, both on the basis of amount of off-target looks and on the basis of where the child was looking at target onset.

For accuracy, we then considered per-child minimums either based on number of valid trials, or based on fraction of valid trials in an administration. Fraction of valid trials is the more common exclusion in the literature, but when drawing from many datasets with different trials counts, minimum trial count may be a cleaner metric. For RT, we considered different cutoffs for minimum numbers of valid trials.

We note that datasets contributed to Peekbank vary in whether excluded data was contributed, or whether only cleaned post-exclusion data was contributed. Thus, some datasets may not show effects of exclusions if those trials or administrations that would be excluded are missing from Peekbank.

Results. For exclusions, there are both considerations around reliability and validity, and around how much of the potentially useable data is being discarded. For trial-level exclusions, a requirement to look at the target and distractor at least 50% of the time results in 85.64 % (range:27.22% - 98.92%) trials being retained. These exclusions lead to an expected increase in ICC by 0.03 [0, 0.05] and no substantial change in test-retest reliability (-0.02 [-0.11, 0.08]). Differences in CDI correlation are not substantive (0.01 [-0.04, 0.06] for comprehension, and 0 [-0.03, 0.03] for production).

Overall, there is a trade-off between reducing the amount of data and increasing the reliability of the remaining data. Filtering based on initial looking position, or having requirements for looking above .5 start to diminish the amount of data, and potentially negatively effect test-retest reliability and CDI measures (Figure S13 and Figure S14). Requirements for on-image looking rates between 0 (no exclusions) and .5 are reasonable.

We then explored the effects of excluding children for too few trials. These are typically set as a fraction, but we also try with total number of trials because the number of trials in different datasets differs.

[could switch out for 50% instead of min 8] Only including administrations with at least 8 trials results in 80.99 % (range:0% - 100%) of admins remaining.

These exclusions lead to an expected increase in ICC by 0 [-0.01, 0.02] and increase in test-retest reliability by 0.03 [-0.08, 0.14]. Differences in CDI correlation are not substantive (0.01 [-0.04, 0.07] for comprehension, and 0.01 [-0.03, 0.05] for production). We do not observe significant differences based on minimum trials below 10 or minimum fractions (Figs S15-S18). It's acceptable, but probably not necessary, to exclude children for too few trials.

We also considered the effects of minimum trial requirements on RT data. Only

including administrations that contribute at least 2 trials, results in 90.93 % (range:42.86% - 100%) of admins remaining.

These exclusions lead to a negligible change in ICC by -0.01 [-0.07, 0.04], a slight increase in test-retest reliability by 0.03 [-0.08, 0.14], no change in CDI comprehension correlation(0 [-0.14, 0.15]), and a possible increased in CDI production correlation (0.02 [-0.04, 0.08]).

For RT, it's reasonable to set low minimums for number of trials (like 2 trials) which does make for stronger correlations with CDI production. There are much stronger tradeoffs between test-retest reliability and number of children included (Figures S19 and S20). It's also reasonable to not exclude.

Overall, exclusions based on amount of data per-trial or per-child come with trade-offs between amount of data and cleanliness of the data. Depending on the study design and goals, no exclusions, or light exclusions can both be reasonable choices.

Design recommendations

Our focus for this work was to explore how data processing choices affected the reliability and validity of the resultant measures. However, along the way, we came across a couple key experimental design questions that could impact reliability and validity.

Recommendation 5: Design studies to have a minimum of 10 trials per child, per condition, or as many trials as possible for individual-difference studies.

Excluding children for low trial counts may increase test-retest reliability, but it also decreases the power and generalizability of results by excluding some children. Rather than retrospective exclusions, another approach is prospective, how many trials per child should one aim for to achieve adequate reliability and validity? Test-retest reliability is particularly important for studies that look at longitudinal effects as it assumes that the test sessions are representative of properties of the child at that age. [something because evidently not everyone who does work that assumes t-rt thinks of it that way]

To address this prospective question, we considered how reliability and validity measures would have changed in we had fewer trials per child. We filtered the data to include only children who had contributed at least N trials (for different values of N) and then repeatedly subsampled their trials and calculated our metrics on the subsets.

The results for accuracy and RT are shown in Figure 6. ICC reliability is not sensitive to the number of trials per subject, so is not shown.

For accuracy, more trials per administration increases the test-retest reliability, up to 20 trials (the highest we measured), although the gains per additional trial decrease with more trials. Validity with CDI measures also increased as trials increased, with smaller gains past 5 trials. Based on accuracy, we recommend at least 10 trials per condition in general, and as many as possible for individual differences studies where test-retest reliability is essential.




Figure 6. Simulations of measurement properties for different trial counts, created by downsampling trials for all children with different trial minimums. Shown for accuracy window of 600 to 4000ms with no trial- or kid- level exclusions.

Nearly every trial a child does results in a useable accuracy measure, but only around 1/3 of trials result in a valid RT measure (in other cases, children start fixating the target or fixating neither).

For test-retest reliability, having at least 5 RT measures per administration is better than having fewer; gains beyond 5 are hard to measure because of low numbers of subjects with repeat administrations and enough trials, but more trials appears to be beneficial. For CDI measures, there are substantial gains for having 2 trials compared to 1, and benefits to additional trials up to 7 trials. For individual differences in processing speed, as many trials as possible is very important (although we are limited by existing datasets to fully quantify the trade-offs).

The validity and test-retest reliability of accuracy and RT increases as the number of trials increases; with large initial gains. When validity matters most, we recommend aims for 10 trials per child per condition, for expected yields of 10 trials for accuracy and 3 for RT. For individual differences when test-retest matters, as many trials as is feasible is recommended. As high as we can simulate based on the available data, there are still gains for more trials.

For this analysis, we treat distinct sessions separately in order to calculate test-retest reliability, but as recommended by Fernald et al. (2008), one way of increasing trial count

especially for individual-difference studies is to have children return for another session and then treat the sessions together (see e.g., Weaver et al., 2024).

Recommendation 6: Match target and distractor images based on animacy.

As mentioned in Recommendation 2, one concern with LWL trial design is that children may have preferences for which image to look at separate from the auditory stimuli, and some of their looking behavior may be independent of the auditory stimulus. One approach is to decrease the preference discrepancy between a pair of images, so that overall, in the absence of auditory cues, children prefer to look at both images approximately equally. While there are a number of factors that could play into children’s image preferences (Fernald et al., 2008; Nothdurft, 2000; Pomper & Saffran, 2019), one factor that stood out in the data was animacy. Children have a strong preference for looking at animals and humans (and parts of humans). When trials have one animate and one inanimate, children look more at the animate in the pre-stimulus period (Figure 7) and are slower to respond when the inanimate is named. In contrast, for trials with target and distractor matched on animacy, children’s early looking patterns are not sensitive to animacy. This difference is particularly relevant for the younger children, whose stimulus-contingent looking behavior is weakest, so more of the looking behavior is driven by animacy; however, effects of animacy were apparent in children up to the age of three (Figure S21).




Figure 7. Looking time profiles for different combinations of target and distractor animacy.

Another potential concern regarding stimulus preferences is that repeating the same pairs of images may cause children to prefer to look at the image that was previously the target [citation – does anyone admit to worrying about this?]. We investigated this in a series of models (reported in SI: Order Effects), and found that effects of repeating target labels (and target and distractor images) are likely quite small, with unclear directionality. In general, our data suggest that order effects stemming from repeated trials of the same target are unlikely to significantly bias word recognition estimates.




Figure 8. Differences in dependent measures between the recommended accuracy measurement and a sensible-sounding, but unfortunate alternative set of parameters.

How much can this matter?

To quantify the potential cumulative consequences of different combinations of choices, we computed the difference in reliability when making all of the recommended choices with respect to data processing (in particular, measuring accuracy on 600-4000 with no exclusions) to an instance in which researchers made a set of reasonable-sounding but unfortunate choices.

DO WE WANT worst case? this is currently a combination of things people actually do (at least individually – 300-1800 bc to -2000 - 0; exclude for .6 and if not looking at start; then require 50% trials)

Switching to the recommended parameters yields an ICC gain of 0.2 [0.13, 0.27] and increases test-retest reliability by 0.21 [0.13, 0.29]. There are also smaller expected increases in validity of (0.07 [-0.03, 0.17] for correlation with comprehension, and 0.07 [0, 0.13] for correlation with production.

For RT, differences are a lot smaller, do we want to include?

Discussion

Children’s early word recognition is an important construct for the study of language learning, but the use of the LWL paradigm is hindered by uncertainty about data analysis. How can we maximize reliable, valid signal in the analysis of data from this paradigm? We took a practical, data-driven approach to this question, using a large-scale dataset of infant looking-while-listening data to quantify the impact of different analysis choices on core measurement properties, leading to six core recommendations for data processing and

experimental design.

For data processing, our strongest take-homes from these analyses concern window choice and baseline correction. First, longer time windows (e.g., windows that end around 3000 ms post target word onset) maximize the reliability of accuracy measures, with no negative consequences for validity. While average accuracy tends to peak earlier in the timecourse of a trial (typically 1000-2000 ms post target word onset), leading shorter windows to be preferred in the literature, these findings suggest that later looking windows still capture signal that boosts the consistency of word recognition measures. Second, baseline correction, while an intuitive solution to imbalances in the saliency of images, causes large decrements in reliability; we therefore generally recommend avoiding this practice, especially when maximizing reliability is a high priority (e.g., when investigating individual differences). Finally, we formulate recommendations for choosing reaction time cutoffs (around 400ms tends to maximize reliability) and find that implementing trial- and child-level exclusions based on the amount of data contributed tend to have minimal consequences. Overall, choices about windows of analysis, approaches to exclusions, and how scores are computed can all have meaningful impacts on measure reliability.

One fascinating question that emerges from our analysis concerns the purity of the construct that LWL measures. In our analysis, long accuracy windows increased reliability more than they did validity. One speculative reason for this difference is that signal from later in the timecourse might not be as tightly coupled to language per se, and hence might instead reflect more general features of attention. Future research could use this difference as a starting point to further parcellate the attentional and linguistic contributors to language-driven eye-movements.

A more general take-home from our data processing analyses is that intuitive data-based practices, e.g., looking at patterns in grand-average time courses, may not always be a trustworthy guide to child- or item-level psychometric properties. Past recommendations for time window selection have sometimes been motivated by inspecting consistent patterns of looking in grand-averages, and, e.g., noting when looking begins to decline across participants. These heuristics — while based in important patterns in looking behavior — may not always lead to optimal choices with respect to core measurement properties such as reliability. Wherever possible, we recommend that measurement choices be guided by direct estimates of psychometric properties in order to maximize the signal in our measures of infant behavior (Byers-Heinlein et al., 2022; Sander-Montant et al., 2025; Weaver et al., 2026).

Among our design-focused recommendations, the most striking result is the impact of animacy on infant looking behavior. When an image of an animate referent was paired with an image of an inanimate referent, infant looking was systematically biased towards the animate image. This finding has important consequences for estimating word recognition: accuracy for animate targets is substantially overestimated and accuracy for inanimate targets becomes substantially underestimated when animacy is not balanced in the target-distractor pairing. The results are especially dramatic for younger infants (under 18

months): looking for familiar inanimate targets can appear to be at chance when paired with an animate distractor. In other words, an improper balancing of saliency could make an infant appear to not know a word, even though they do. Once animacy is balanced in the design, animate and inanimate targets have highly similar recognition curves across all ages. These results emphasize the importance of balancing saliency in image pairings as much as possible in looking-while-listening designs. This is especially true given that many of the current analytic approaches to correcting for saliency imbalances *post-hoc* rely on difference-based measures (baseline correction; item-pair correction) which decrease reliability.

Limitations & Generalizability

The scope of our recommendations is fundamentally limited by the characteristics of the Peekbank dataset. First, while the Peekbank database is by far the largest dataset of its kind, and one of the largest harmonized infant eyetracking datasets, it is not a comprehensive or representative set of looking-while-listening studies. To highlight just a few of these limitations, (1) the data is predominantly English word recognition collected from infants living in the U.S., and (2) the dataset is a convenience sample from labs who made their data openly available or otherwise accessible on request from the Peekbank team. Looking-while-listening studies as a whole are likely more heterogeneous than what is represented in Peekbank — the datasets considered here come predominantly from a small group of labs or networks of researchers with similar training in looking-while-listening data collection.

At the same time, there is also substantial heterogeneity in the studies included in Peekbank, due to many studies pursuing different research questions and therefore making a variety of distinct design choices (e.g., surrounding target words, image properties, and characteristics of the auditory stimuli). To mitigate this concern, our analyses focused exclusively on what we term “vanilla” trials, i.e., trials that present familiar words with familiar images with limited or no targeted manipulations of the stimuli. The main benefit of focusing on vanilla trials is that it allows for a clear interpretation of reliability and validity — trials are truly designed to measure word recognition, rather than some other aspect of language processing. We view the remaining heterogeneity of study properties (e.g., specific words; differences in age) as a strength of our approach — our findings generalize broadly across a wide range of items and ages. However, by focusing on vanilla trials only, we also limit our conclusions to measuring familiar word recognition; our findings do not address how data processing and design practices would generalize to other estimation goals, e.g., estimating the magnitude of differences between two experimental conditions (though there is good reason to think that some of the more consequential choices, e.g. baseline correction, are likely have a similar impact).

While we summarize the choices that typically maximize reliability, we note that we often see a range of choices with similar outcomes in our analyses. For example, in determining the best time window for computing accuracy, a choice of ending a time window at 2800 vs. 3000 or 3000 ms vs. 3200 ms leads to very similar outcomes. There is no one choice of time window that we can definitively infer as the “best” for maximizing reliability —

instead, we see a pattern of graded improvements in reliability that follow a predictable pattern (longer time windows tend to increase reliability holding the onset of the accuracy window constant). While we sometimes articulate which choices lead to the highest reliability descriptively, these are at best approximations. We offer these windows as data-driven default choices (especially in cases where the decision would otherwise be ad-hoc).

Conclusions

Experimental design and data processing decisions can have profound impacts on the properties of infant measures. Our findings show how compiling large-scale datasets such as Peekbank can help us estimate the impact of a wide variety of experimenter decisions, generalizable across a wide range of ages, stimuli, and experimental contexts. More generally, establishing consistent practices and principles for reliable and valid measurement of infant looking sets the stage for asking central questions about how language skill develops.

References

- Adams, K. A., Marchman, V. A., Loi, E. C., Ashland, M. D., Fernald, A., & Feldman, H. M. (2018). Caregiver talk and medical risk as predictors of language outcomes in full term and preterm toddlers. *Child Development*, 89(5), 1674–1690.
- Bacon, D., & Saffran, J. (2022). Role of speaker gender in toddler lexical processing. *Infancy : The Official Journal of the International Society on Infant Studies*, 27(2), 291–300.
- Bergelson, E., & Swingley, D. (2012). At 6-9 months, human infants know the meanings of many common nouns. *Proceedings of the National Academy of Sciences*, 109(9), 3253–3258.
- Bergmann, C., Tsuji, S., Piccinini, P. E., Lewis, M. L., Braginsky, M., Frank, M. C., & Cristia, A. (2018). Promoting replicability in developmental research through meta-analyses: Insights from language acquisition research. *Child Development*, 89(6), 1996–2009.
- Borovsky, A., & Peters, R. E. (2019). Vocabulary size and structure affects real-time lexical recognition in 18-month-olds. *PLOS ONE*, 14(7), e0219290. <https://doi.org/10.1371/journal.pone.0219290>
- Byers-Heinlein, K., Bergmann, C., & Savalei, V. (2022). Six solutions for more reliable infant research. *Infant and Child Development*, 31(5), e2296. <https://doi.org/10.1002/icd.2296>
- Csibra, G., Hernik, M., Mascaró, O., Tatone, D., & Lengyel, M. (2016). Statistical treatment of looking-time data. *Developmental Psychology*, 52(4), 521–536.
- Fenson, L., Dale, P. S., Reznick, J. S., Bates, E., Thal, D. J., Pethick, S. J., . . . Stiles, J. (1994). Variability in early communicative development. *Monographs of the Society for Research in Child Development*, i–185.
- Fernald, A., & Marchman, V. A. (2012). Individual differences in lexical processing at 18 months predict vocabulary growth in typically developing and late-talking toddlers. *Child Development*, 83(1), 203–222.
- Fernald, A., Perfors, A., & Marchman, V. A. (2006). Picking up speed in understanding: Speech processing efficiency and vocabulary growth across the 2nd year. *Developmental Psychology*, 42(1), 98–116. <https://doi.org/10.1037/0012-1649.42.1.98>
- Fernald, A., Pinto, J. P., Swingley, D., Weinberg, A., & McRoberts, G. W. (1998). Rapid gains in speed of verbal processing by infants in the 2nd year. *Psychological Science*, 9(3), 228–231.
- Fernald, A., Zangl, R., Portillo, A. L., & Marchman, V. A. (2008). Looking while listening: Using eye movements to monitor spoken language comprehension by infants and young children. In I. A. Sekerina, E. M. Fernández, & H. Clahsen (Eds.), *Language Acquisition and Language Disorders* (Vol. 44, pp. 97–135). Amsterdam: John Benjamins Publishing Company. <https://doi.org/10.1075/lald.44.06fer>
- Flake, J. K., & Fried, E. I. (2020). Measurement Schmeasurement: Questionable Measurement Practices and How to Avoid Them. *Advances in Methods and Practices in Psychological Science*, 3(4), 456–465. <https://doi.org/10.1177/2515245920952393>
- Frank, M. C., Braginsky, M., Yurovsky, D., & Marchman, V. A. (2017). Wordbank: An open repository for developmental vocabulary data. *Journal of Child Language*, 44(3), 677–694.
- Frank, M. C., Braginsky, M., Yurovsky, D., & Marchman, V. A. (2021). *Variability and*

- consistency in early language learning: *The wordbank project*. Cambridge, MA: MIT Press.
- Frank, M. C., Marchman, V. A., Bergey, C. A., Boyce, V., Braginsky, M., Kachergis, G., ... Zettersten, M. (2025). Continuous developmental changes in word recognition support language learning across early childhood. *eLife*. <https://doi.org/10.7554/eLife.109636.1>
- Garrison, H., Baudet, G., Breitfeld, E., Aberman, A., & Bergelson, E. (2020). Familiarity plays a small role in noun comprehension at 12–18 months. *Infancy : The Official Journal of the International Society on Infant Studies*, 25(4), 458–477.
- Golinkoff, R. M., Ma, W., Song, L., & Hirsh-Pasek, K. (2013). Twenty-five years using the intermodal preferential looking paradigm to study language acquisition: What have we learned? *Perspectives on Psychological Science*, 8(3), 316–339.
- Hirsh-Pasek, K., & Golinkoff, R. M. (1996). The intermodal preferential looking paradigm: A window onto emerging language comprehension. In *Language, Speech, and Communication. Methods for assessing children's syntax* (pp. 105–124). Cambridge, MA, US: The MIT Press.
- Hurtado, N., Marchman, V. A., & Fernald, A. (2007). Spoken word recognition by Latino children learning Spanish as their first language. *Journal of Child Language*, 34(2), 227–249. <https://doi.org/10.1017/S0305000906007896>
- Kominsky, J. F., Lucca, K., Thomas, A. J., Frank, M. C., & Hamlin, J. K. (2022). Simplicity and validity in infant research. *Cognitive Development*, 63, 101213. <https://doi.org/10.1016/j.cogdev.2022.101213>
- Law, F., II, & Edwards, J. R. (2015). Effects of Vocabulary Size on Online Lexical Processing by Preschoolers. *Language Learning and Development*, 11(4), 331–355. <https://doi.org/10.1080/15475441.2014.961066>
- López Pérez, M., Moore, C., Sander-Montant, A., & Byers-Heinlein, K. (2024). Infants' Knowledge of Individual Words: Investigating Links Between Parent Report and Looking Time. *Infancy*, 30(1), e12641. <https://doi.org/10.1111/inf.12641>
- Lord, F. M. (1956). The Measurement of Growth. *Educational and Psychological Measurement*, 16(4), 421–437. <https://doi.org/10.1177/001316445601600401>
- Mahr, T., McMillan, B. T. M., Saffran, J. R., Ellis Weismer, S., & Edwards, J. (2015). Anticipatory coarticulation facilitates word recognition in toddlers. *Cognition*, 142, 345–350. <https://doi.org/10.1016/j.cognition.2015.05.009>
- Marchman, Virginia A., Adams, K. A., Loi, E. C., Fernald, A., & Feldman, H. M. (2016). Early language processing efficiency predicts later receptive vocabulary outcomes in children born preterm. *Child Neuropsychology*, 22(6), 649–665. <https://doi.org/10.1080/09297049.2015.1038987>
- Marchman, V. A., Dale, P. S., & Fenson, L. (2023). *MacArthur-bates communicative development inventories users' guide and technical manual, third edition*. Brookes.
- Marchman, Virginia A., Loi, E. C., Adams, K. A., Ashland, M., Fernald, A., & Feldman, H. M. (2018). Speed of language comprehension at 18 months old predicts school-relevant outcomes at 54 months old in children born preterm. *Journal of Developmental & Behavioral Pediatrics*, 39(3), 246–253.
- Mirman, D., Dixon, J. A., & Magnuson, J. S. (2008). Statistical and computational models of the visual world paradigm: Growth curves and individual differences. *Journal of Memory and Language*, 59(4), 475–494. <https://doi.org/10.1016/j.jml.2007.11.006>

- Moore, C., & Bergelson, E. (2024). Wordform variability in infants' language environment and its effects on early word learning. *Cognition*, *245*, 105694. <https://doi.org/10.1016/j.cognition.2023.105694>
- Nothdurft, H.-C. (2000). Salience from feature contrast: Additivity across dimensions. *Vision Research*, *40*(10-12), 1183–1201.
- Novack, M. A., Dworak, E. M., Han, Y. C., Kaat, A. J., Ustsinovich, V., Saffran, J., . . . Gershon, R. C. (2025). Development and validation of the NIH Baby Toolbox® language measures. *Infant Behavior and Development*, *80*, 102121. <https://doi.org/10.1016/j.infbeh.2025.102121>
- Olson, R. H., Pomper, R., Potter, C. E., Hay, J. F., Saffran, J. R., Ellis Weismer, S., & Lew-Williams, C. (n.d.). *Peyecoder: An open-source program for coding eye movements*. <https://doi.org/10.5281/zenodo.3939234>
- Peelle, J. E., & Van Engen, K. J. (2021). Time Stand Still: Effects of Temporal Window Selection on Eye Tracking Analysis. *Collabra: Psychology*, *7*(1), 25961. <https://doi.org/10.1525/collabra.25961>
- Perry, L. K., & Saffran, J. R. (2017). Is a pink cow still a cow? Individual differences in toddlers' vocabulary knowledge and lexical representations. *Cognitive Science*, *41*(4), 1090–1105. <https://doi.org/10.1111/cogs.12370>
- Peter, M. S., Durrant, S., Jessop, A., Bidgood, A., Pine, J. M., & Rowland, C. F. (2019). Does speed of processing or vocabulary size predict later language growth in toddlers? *Cognitive Psychology*, *115*, 101238.
- Pomper, R., & Saffran, J. R. (2016). Roses are red, socks are blue: Switching dimensions disrupts young children's language comprehension. *PloS One*, *11*(6), e0158459. <https://doi.org/10.1371/journal.pone.0158459>
- Pomper, R., & Saffran, J. R. (2019). Familiar object salience affects novel word learning. *Child Development*, *90*(2), e246–e262. <https://doi.org/10.1111/cdev.13053>
- Potter, C. E., Fourakis, E., Morin-Lessard, E., Byers-Heinlein, K., & Lew-Williams, C. (2019). Bilingual toddlers' comprehension of mixed sentences is asymmetrical across their two languages. *Developmental Science*, *22*(4), e12794. <https://doi.org/10.1111/desc.12794>
- Potter, C., & Lew-Williams, C. (unpublished). *Behold the canine!: How does toddlers' knowledge of typical frames and familiar words interact to influence their sentence processing?*
- Reznick, J. S. (1990). Visual preference as a test of infant word comprehension. *Applied Psycholinguistics*, *11*(2), 145–166.
- Rhemtulla, M., & Bottesini, J. G. (2022). On the reliability of reliability estimates in infancy research. *Infant and Child Development*, *31*(5), e2327. <https://doi.org/10.1002/icd.2327>
- Rogosa, D. R., & Willett, J. B. (1983). Demonstrating the Reliability the Difference Score in the Measurement of Change. *Journal of Educational Measurement*, *20*(4), 335–343. <https://doi.org/10.1111/j.1745-3984.1983.tb00211.x>
- Ronfard, S., Wei, R., & Rowe, M. L. (2022). Exploring the linguistic, cognitive, and social skills underlying lexical processing efficiency as measured by the looking-while-listening paradigm. *Journal of Child Language*, *49*(2), 302–325. <https://doi.org/10.1017/S0305000921000106>
- Sander-Montant, A., Fibla, L., Byers-Heinlein, K., & Killam, H. (2025). *Keeping an eye on*

- looking measures: Towards more robust developmental methods. OSF.
https://doi.org/10.31234/osf.io/vxu53_v1
- Sander-Montant, A., Pérez, M. L., & Byers-Heinlein, K. (2023). The more they hear the more they learn? Using data from bilinguals to test models of early lexical development. *Cognition*, 238, 105525. <https://doi.org/10.1016/j.cognition.2023.105525>
- Schreiner, M. S., Zettersten, M., Bergmann, C., Frank, M. C., Fritzsche, T., Gonzalez-Gomez, N., et al.others. (2024). Limited evidence of test-retest reliability in infant-directed speech preference in a large preregistered infant experiment. *Developmental Science*, 27(6), e13551. <https://doi.org/10.1111/desc.13551>
- Snedeker, J. (2013). Children's sentence processing. In *Sentence Processing*. Psychology Press.
- Swingley, D., & Aslin, R. N. (2000). Spoken word recognition and lexical representation in very young children. *Cognition*, 76(2), 147–166.
[https://doi.org/10.1016/s0010-0277\(00\)00081-0](https://doi.org/10.1016/s0010-0277(00)00081-0)
- Swingley, D., & Aslin, R. N. (2002). Lexical Neighborhoods and the Word-Form Representations of 14-Month-Olds. *Psychological Science*, 13(5), 480–484.
<https://doi.org/10.1111/1467-9280.00485>
- Trafimow, D. (2015). A defense against the alleged unreliability of difference scores. *Cogent Mathematics*, 2(1), 1064626. <https://doi.org/10.1080/23311835.2015.1064626>
- Venker, C. E., Pomper, R., Mahr, T., Edwards, J., Saffran, J., & Ellis Weismer, S. (2020). Comparing automatic eye tracking and manual gaze coding methods in young children with autism spectrum disorder. *Autism Research*, 13(2), 271–283.
<https://doi.org/https://doi.org/10.1002/aur.2225>
- Von Holzen, K., & Mani, N. (2012). Language nonselective lexical access in bilingual toddlers. *Journal of Experimental Child Psychology*, 113(4), 569–586.
<https://doi.org/10.1016/j.jecp.2012.08.001>
- Wass, S. V., Smith, T. J., & Johnson, M. H. (2013). Parsing eye-tracking data of variable quality to provide accurate fixation duration estimates in infants and adults. *Behavior Research Methods*, 45(1), 229–250. <https://doi.org/10.3758/s13428-012-0245-6>
- Weaver, H., & Saffran, J. R. (2026). Interrogating Early Word Knowledge: Factors That Influence the Alignment Between Caregiver-Report and Experimental Measures. *Developmental Science*, 29(1), e70088. <https://doi.org/10.1111/desc.70088>
- Weaver, H., Saffran, J., & Arunachalam, S. (2026). Looking forward: Eye-gaze methods in vocabulary development research. *Trends in Cognitive Sciences*.
- Weaver, H., Zettersten, M., & Saffran, J. R. (2024). Becoming word meaning experts: Infants' processing of familiar words in the context of typical and atypical exemplars. *Child Development*, 95(5), e352–e372. <https://doi.org/10.1111/cdev.14120>
- Weisleder, A., & Fernald, A. (2013). Talking to children matters: Early language experience strengthens processing and builds vocabulary. *Psychological Science*, 24(11), 2143–2152.
<https://doi.org/10.1177/0956797613488145>
- Yurovsky, D., Wade, A., Kraus, A., Gengoux, G., Antonio, Y., & Frank, M. (2017). *Developmental changes in the speed of social attention in early word learning*. unpublished manuscript.
- Zettersten, M., Yurovsky, D., Xu, T. L., Uner, S., Tsui, A. S. M., Schneider, R. M., ... Frank, M. C. (2022). Peekbank: An open, large-scale repository for developmental

812 eye-tracking data of children's word recognition. *Behavior Research Methods*, 55(5),
813 2485–2500. <https://doi.org/10.3758/s13428-022-01906-4>